

Unit 1: Overview of Operating Systems

1.3. Windows Operating System Family - Concepts & Tools

Roadmap for Section 1.3.

High-level Overview on Windows Concepts

- Processes, Threads
- Virtual Memory, Protection
- Objects and Handles

Windows is thoroughly instrumented

- Key monitoring tools
- Extra resources at www.sysinternals.com

Requirements and Design Goals for the original Windows NT project

- Provide a true 32-bit, preemptive, reentrant, virtual memory operating system
- Run on multiple hardware architectures and platforms
- Run and scale well on symmetric multiprocessing systems
- Be a great distributed computing platform (Client & Server)
- Run most existing 16-bit MS-DOS and Microsoft Windows 3.1 applications
- Meet government requirements for POSIX 1003.1 compliance
- Meet government and industry requirements for operating system security
- Be easily adaptable to the global market by supporting Unicode

Goals (contd.)

- **Extensibility**

- Code must be able to grow and change as market requirements change.

- **Portability**

- The system must be able to run on multiple hardware architectures and must be able to move with relative ease to new ones as market demands dictate.

- **Reliability and Robustness**

- Protection against internal malfunction and external tampering.
- Applications should not be able to harm the OS or other running applications.

- **Compatibility**

- User interface and APIs should be compatible with older versions of Windows as well as older operating systems such as MS-DOS.
- It should also interoperate well with UNIX, OS/2, and NetWare.

- **Performance**

- Within the constraints of the other design goals, the system should be as fast and responsive as possible on each hardware platform.

Portability

- HAL (Hardware Abstraction Layer):
 - support for x86 (initial), MIPS (initial), Alpha AXP, PowerPC (NT 3.51), Itanium (Windows XP/2003)
 - Machine-specific functions located in HAL
- Layered design:
 - architecture-specific functions located in kernel
- Windows kernel components are primarily written in C:
 - OS executive, utilities, drivers
 - UI and graphics subsystem - written in C++
 - HW-specific/performance-sensitive parts - written in assembly language: interrupt trap handler, context switching

Windows API & Subsystems

- Windows API (application programming interface):
 - Common programming interface to three different Windows OSes: Windows NT, Windows 9x and Windows CE (Embedded Compact)
 - OSes implement (different) subsets of the API
 - MSDN: <http://msdn.microsoft.com>
- Windows supports multiple subsystems (APIs):
 - Windows (primary), POSIX, OS/2
 - User space application access OS functionality via subsystems
- Subsystems define APIs, process, and file system semantics
 - OS/2 used to be primary subsystem for Windows NT 3.x

64-bit vs. 32-bit Windows APIs

- Pointers and types derived from pointer, e.g. handles, are 64-bit long
 - A few others go 64, e.g. WPARAM, LPARAM, LRESULT, SIZE_T
 - Rest are the same, e.g., 32-bit INT, DWORD, LONG

Win32 and Win64
are consistently
named the
Windows API

API	Data Model	<code>int</code>	<code>long</code>	pointer
Win32	ILP32	32	32	32
Win64	LLP64	32	32	64
UNIXes	LP64	32	64	64

Services, Functions, and Routines

- Windows API functions:
 - Documented, callable subroutines
 - *CreateProcess*, *CreateFile*, *GetMessage*
- Windows system services:
 - Undocumented functions, callable from user space
 - *NtCreateProcess* is used by Windows *CreateProcess* and POSIX *fork()* as an internal service
- Windows internal routines:
 - Subroutines inside the Windows executive, kernel, or HAL
 - Callable from kernel mode only (device driver, NT OS components)
 - *ExAllocatePool* allocates memory on Windows system heap

Services, Functions, and Routines (contd.)

- Windows services:

- Processes which are started by the Service Control Manager
- Example: The *Schedule* service supports the at-command

- DLL (dynamic link library)

- Subroutines in binary format contained in dynamically loadable files
- Examples: MSVCRT.DLL – MS Visual C++ run-time library
 KERNEL32.DLL – one of the Windows API libraries

Processes and Threads

What is a process?


- Represents an instance of a running program
 - you create a process to run a program
 - starting an application creates a process
- Process defined by:
 - Address space
 - Resources (e.g. open handles)
 - Security profile (token)

What is a thread?

- An execution context within a process
- Unit of scheduling (threads run, processes don't run)
- All threads in a process share the same per-process address space
 - Services provided so that threads can synchronize access to shared resources (critical sections, mutexes, events, semaphores)
- All threads in the system are scheduled as peers to all others, without regard to their "parent" process

System calls

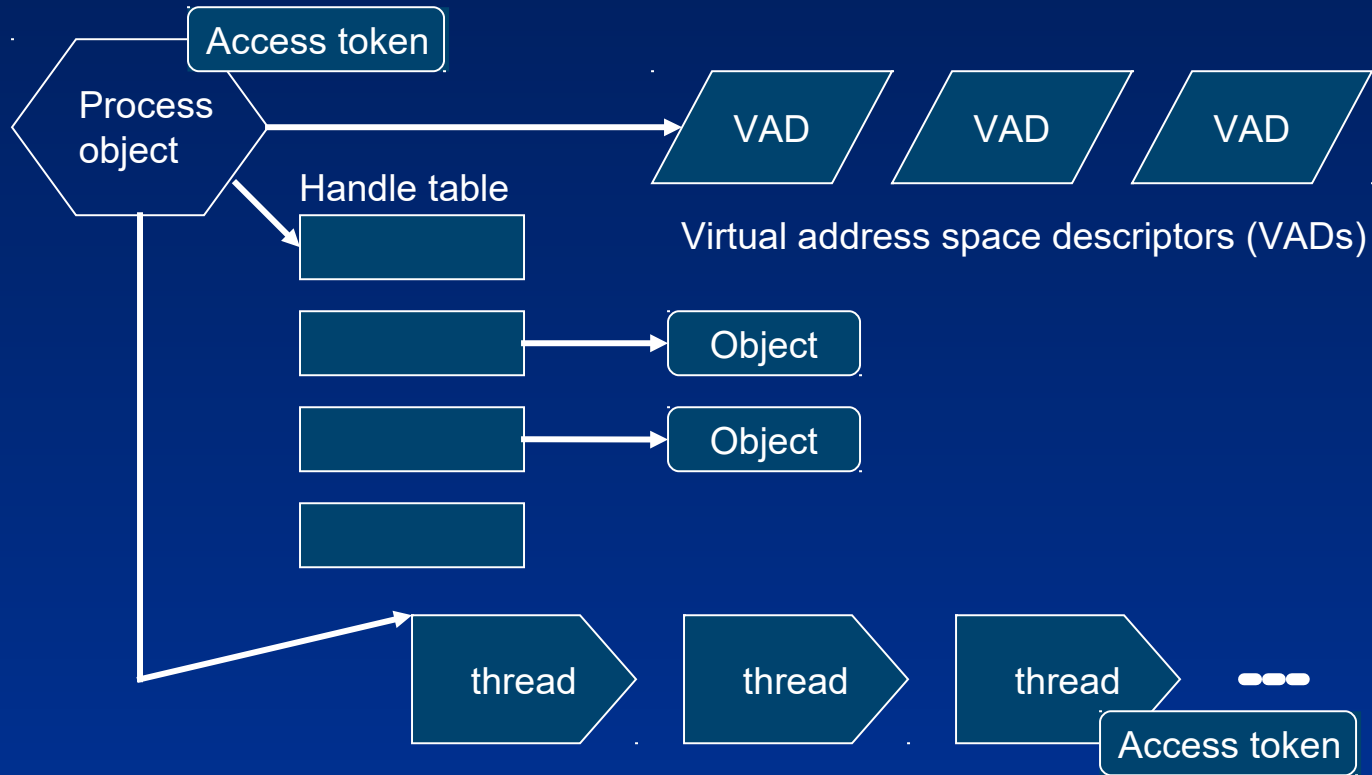
- Primary argument to CreateProcess is image file name (or command line)
- Primary argument to CreateThread is a function entry point address



Processes & Threads

- Every process starts with one thread
 - First thread executes the program's "main" function
 - Can create other threads in the same process
 - Can create additional processes
- Why divide an application into multiple threads?
 - Perceived user responsiveness, parallel/background execution
 - Examples: Word background print – can continue to edit during print
 - Take advantage of multiple processors
 - On an MP system with n CPUs, n threads can literally run at the same time
 - Question: given a single threaded application, will adding a 2nd processor make it run faster?
 - Does add complexity
 - Synchronization
 - Scalability well is a different question...
 - # of multiple runnable threads vs # CPUs
 - Having too many runnable threads causes excess context switching

A Process and its Resources



Virtual Memory


- 32-bit address space (4 GB)
 - 2 GB user space (per process)
 - 2 GB operating system
- 64-bit address space
 - 7192 GB user space (Itanium)
 - 8192 GB user space (x64)
 - ~6000 GB operating system
- Memory manager maps virtual onto physical memory

Default 32-bit layout

Unique per process



Systemwide



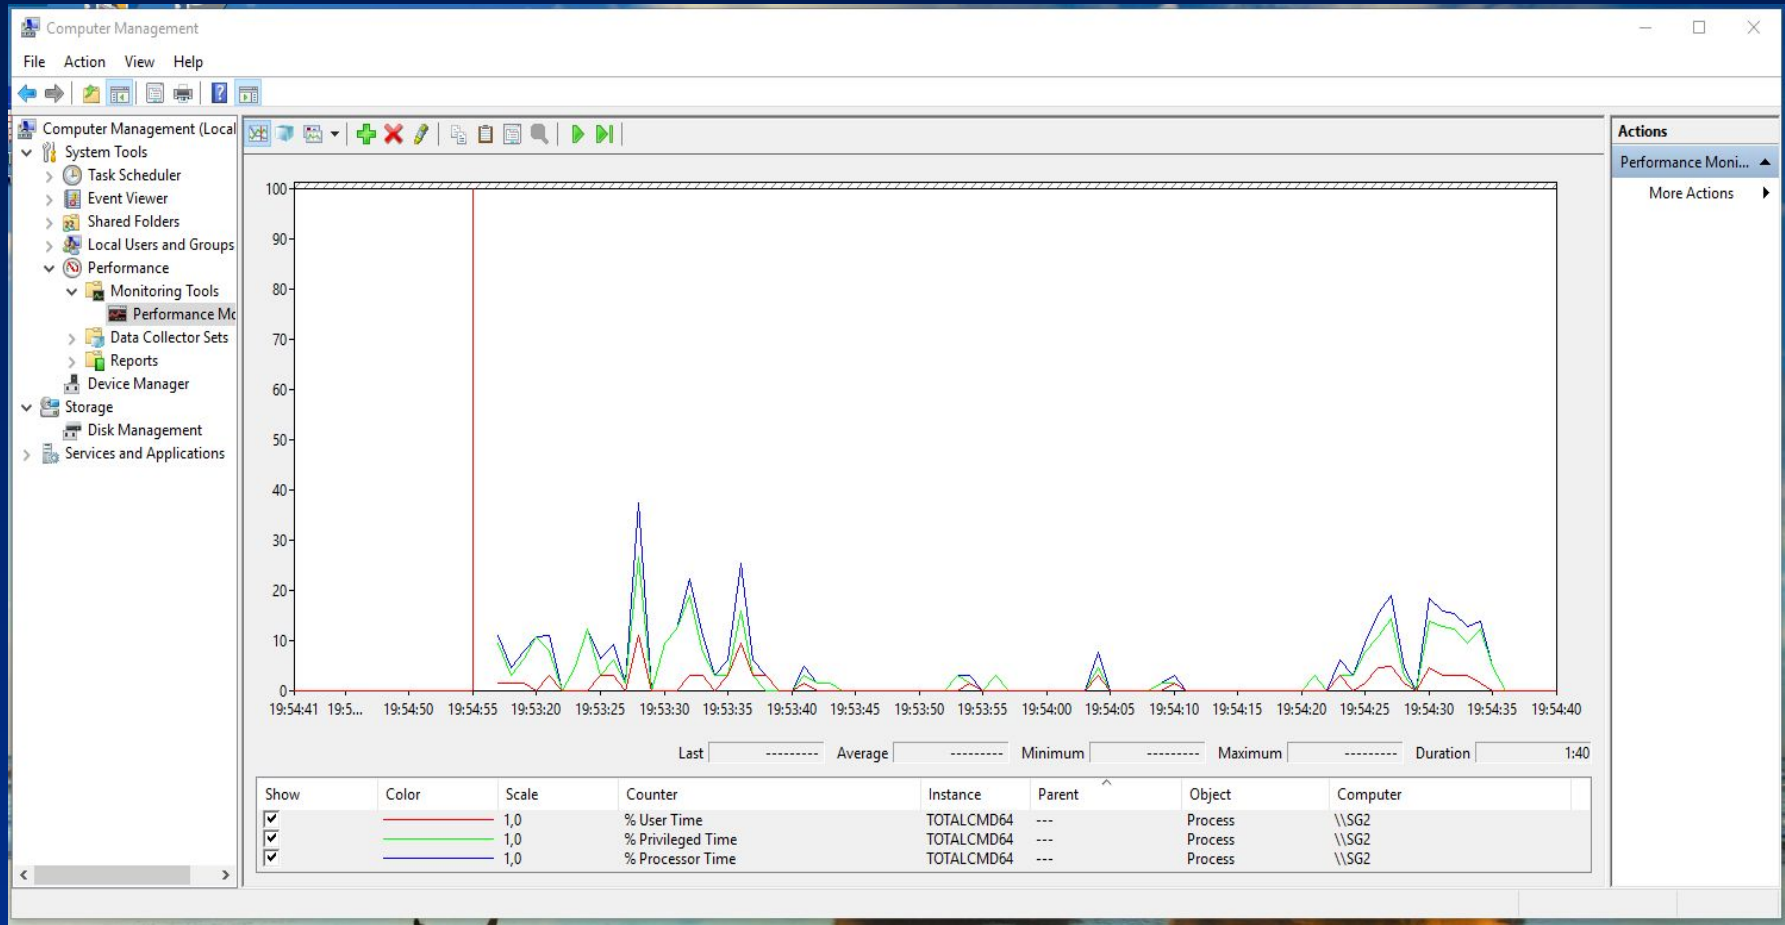
Memory Protection Model

- No user process can touch another user process address space (without first opening a handle to the process, which means passing through NT security)
 - Separate process page tables prevent this
 - “Current” page table changed on context switch from a thread in 1 process to a thread in another process
- No user process can touch kernel memory
 - Page protection in process page tables prevent this
 - OS pages only accessible from “kernel mode”
 - Threads change from user to kernel mode and back (via a secure interface) to execute kernel code
 - Does not affect scheduling (not a context switch)

Kernel Mode vs. User Mode

- No protection for components running in kernel mode
- Transition from user mode to kernel mode through special instruction (processor changes privilege level)
 - OS traps this instruction and validates arguments to syscalls
 - Transition from user to kernel mode does not affect thread scheduling
- Performance Counters: System/Processor/Process/Thread – Privileged Time/User time
 - Windows kernel is thoroughly instrumented
 - Hundreds of performance counters throughout the system
- Performance Monitor – perfmon.msc – a Microsoft Management Console (MMC) snap in

Performance Monitor



Objects and Handles

- Process, thread, file, event objects in Windows - are mapped on NT executive objects
- Object services read/write object attributes
- Objects:
 - Human-readable names for system resources
 - Resource sharing among processes
 - Resource protection against unauthorized access
- Security/Protection based on NT executive objects
- 2 forms of access control:
 - Discretionary control: read/write/access rights
 - Privileged access: administrator may take ownership of files

Networking

- Integral, application-transparent networking services
 - Basic file and print sharing and using services
- A platform for distributed applications
 - Application-level inter-process communication (IPC)
- Windows provides an expandable platform for other network components

Registry

- System wide software settings: boot & configuration info
- Security database
- Per-user profile settings
- In-memory volatile data (current hardware state)
 - What devices are loaded?
 - Resources used by devices
 - Performance counters are accessed through registry functions
- Regedit.exe is the tool to view/modify registry settings
 - HKEY_LOCAL_MACHINE\System\CurrentControlSet\Control
 - HKEY_LOCAL_MACHINE\System\CurrentControlSet\Services
 - HKEY_LOCAL_MACHINE\Software

Unicode

- Most internal text strings are stored/processed as 16-bit wide Unicode strings
 - Windows API string functions have 2 versions
 - Unicode (wide) version
 - L“This string uses 16-bit characters”
 - ANSI (narrow) version
 - “This string uses 8-bit characters”
 - Generic character representation in Windows API
 - _T (“This string uses generic characters”)
- (Windows 9x has Windows API but no Unicode characters,
Windows CE has Windows API but Unicode characters only)

Tools used to dig in

- Many tools available to dig into Windows internals
 - Helps to see internals behavior “in action”
- Several sources of tools
 - Support Tools
 - Resource Kit Tools
 - Debugging Tools
 - Sysinternals.com
- Additional tool packages with internals information
 - Platform Software Development Kit (SDK)
 - Device Driver Development Kit (DDK)

Tools for Viewing Windows Internals

Tool	Image Name	Origin
Startup Programs Viewer	AUTORUNS	www.sysinternals.com
Dependency Walker	DEPENDS	Support Tools, Platform SDK
DLL List	LISTDLLS	www.sysinternals.com
EFS Information Dumper	EFSDUMP	www.sysinternals.com *
File Monitor	FILEMON	www.sysinternals.com
Global Flags	GFLAGS	Support Tools
Handle Viewer	HANDLE	www.sysinternals.com
Junction tool	JUNCTION	www.sysinternals.com
Kernel debuggers	WINDBG, KD	Debugging tools, Platform SDK, Windows DDK
Live Kernel Debugging	LIVEKD	www.sysinternals.com
Logon Sessions	LOGINSESSIONS	www.sysinternals.com
Object Viewer	WINOBJ	www.sysinternals.com
Open Handles	OH	Resource kits
Page Fault Monitor	PFMON	Support Tools, Resource kits, Platform SDK
Pending File Moves	PENDMOVES	www.sysinternals.com

Tools for Viewing Windows Internals (contd.)

Tool	Image Name	Origin
Performance tool	PERFMON.MSC	Windows built-in tool
PipeList tool	PIPELIST	www.sysinternals.com
Pool Monitor	POOLMON	Support Tools, Windows DDK
Process Explorer	PROCEXP	www.sysinternals.com
Get SID tool	PSGETSID	www.sysinternals.com
Process Statistics	PSTAT	Support Tools, Windows 2000 Resource kits, Platform SDK, www.reskit.com
Process Viewer	PVIEWER (in the Support Tools) or PVIEW (in the Platform SDK)	Platform SDK
Quick Slice	QSLICE	Windows 2000 resource kits
Registry Monitor	REGMON	www.sysinternals.com
Service Control	SC	Windows XP, Platform SDK, Windows 2000 resource kits
Task (Process) List	TLIST	Debugging tools
Task Manager	TASKMGR	Windows built-in tool
TDImon	TDIMON	www.sysinternals.com

Windows Debugging Tools

- Separate package of advanced debugging tools
- Download latest version from:
 - <http://www.microsoft.com/whdc/ddk/debugging>
- Tools
 - User-mode and kernel-mode debuggers
 - Kd – command line interface
 - WinDbg – GUI interface (kernel debugging still mostly “command line”)
 - Allow exploring internal system state & data structures
 - Nttd, Cdb – command line user-mode debugger (newer versions than what ships with OS)
 - Misc other tools (some are also in Support Tools):
 - kill, remote, tlist, logger/logview (API logging tool), Autodump

Live Kernel Debugging

- Useful for investigating internal system state not available from other tools
 - Previously, required 2 computers (host and target)
 - Target would be halted while host debugger in use
- Starting from version XP & Srv2003, Windows NT supports live local kernel debugging
 - Technically requires system to be booted /DEBUG to work correctly
 - You can edit kernel memory on the live system (!)
 - But, not all commands work
- LiveKd (www.sysinternals.com)
 - Tricks standard Microsoft kernel debuggers into thinking they are looking at a crash dump
 - Commands that fail in local kernel debugging work in LiveKD:
 - Kernel stacks (!process, !thread)
 - Lm (list modules)
 - Can snapshot a live system (.dump)
 - Does not guarantee consistent view of system memory
 - Thus can loop or fail with access violation
 - Just quit and restart

Sysinternals Tools

- **Freeware Windows internals tools from www.sysinternals.com**
 - Written by Mark Russinovich & Bryce Cogswell (cofounders of Winternals)
- **Useful for developers, system administrators, and power users**
 - Most popular: Filemon, Regmon, Process Explorer
- **Require no installation – run them directly after downloading and unzipping**
- **Many tools require administrative privileges**
 - Some load a device driver
- **Tools regularly updated, so make sure to check for updated versions**

Process Explorer (Sysinternals)

“Super Task Manager”

- Shows full image path, command line, environment variables, parent process, security access token, open handles, loaded DLLs & mapped files

The screenshot displays the Process Explorer window from Sysinternals. The top pane shows a tree view of processes, with 'winlogon.exe' selected. The bottom pane shows the open handles for the selected process.

Process	PID	CPU	Description	Owner	Session	Ha
System Idle Process	0	0		<access denied>	0	0
System	4	0		NT AUTHORITY...	0	455
smss.exe	396	0	Windows NT Session Manager	NT AUTHORITY...	0	21
csrss.exe	452	0	Client Server Runtime Process	NT AUTHORITY...	0	510
winlogon.exe	476	0	Windows NT Logon Application	NT AUTHORITY...	0	568
explorer.exe	312	0	Windows Explorer	DSOLOMON\ds...	0	679
OUTLOOK.EXE	1312	0	Microsoft Outlook	DSOLOMON\ds...	0	435
cmd.exe	1980	0	Windows Command Processor	DSOLOMON\ds...	0	48
hh.exe	1316	0	Microsoft® HTML Help Executable	DSOLOMON\ds...	0	180
procexp.exe	2932	0	Sysinternals Process Explorer	DSOLOMON\ds...	0	57

Handle	Type	Access	Name
0x634	Desktop	0x000F01FF	\Default
0xAC	Desktop	0x000F01FF	\Winlogon
0xB4	Desktop	0x000F01FF	\Disconnect
0xB8	Desktop	0x000F01FF	\Default
0x14	Directory	0x000F000F	\Windows
0x28	Directory	0x0002000F	\BaseNamedObjects
0x8	Directory	0x00000003	\KnownDlls
0x188	Event	0x00100000	\BaseNamedObjects\WinSta0_DesktopSwitch
0x1A0	Event	0x001F0003	\BaseNamedObjects\ThemesStartEvent
0x1B0	Event	0x001F0003	\BaseNamedObjects\WFP_IDLE_TRIGGER

winlogon.exe pid: 476 Refresh Rate: Paused

Platform SDK

(Software Development Kit)

- **Contains header files, libraries, documentation, & sample code for entire Windows “platform” API**
 - 14 separate SDKs
 - “Core SDK” contains core services, COM, messaging, active directory, management, etc.
- **Freely downloadable from**
www.microsoft.com/msdownload/platformsdk/sdkupdate
 - Part of MSDN Professional (or higher) subscription
- **Always matches operating system revision**
 - E.g. Platform SDK revised with new release (or beta) as new APIs are added
- **Not absolutely required for Win32 development (because VC++ comes with the Win32 API header files), but...**
 - VC++ headers, libs, doc won't reflect APIs added after VC++ was mastered
- **Also provides a few tools (e.g. WinObj, Working Set Tuner) not available elsewhere**

Further Reading

- Pavel Yosifovich, Alex Ionescu, et al., “Windows Internals”, 7th Edition, Microsoft Press, 2017.
 - Concepts and Tools (from chapter 1)
 - Digging into Windows Internals (from chapter 1)