

# Unit 3: Concurrency

## 3.1. Concurrency, Critical Sections, Semaphores

# Roadmap for Section 3.1.

- The Critical-Section Problem
- Software Solutions
- Synchronization Hardware
- Semaphores
- Synchronization in Windows & Linux

# The Critical-Section Problem

- $n$  threads all competing to use a shared resource (i.e.; shared data)
- Each thread has a code segment, called *critical section*, in which the shared data is accessed
- Problem:  
Ensure that when one thread is executing in its critical section, no other thread is allowed to execute in its critical section

# Solution to Critical-Section Problem

## 1. Mutual Exclusion

- Only one thread at a time is allowed into its critical section, among all threads that have critical sections for the same resource or shared data.
- A thread halted in its non-critical section must not interfere with other threads.

## 2. Progress

- A thread remains inside its critical section for a finite time only.
- No assumptions concerning relative speed of the threads.

## 3. Bounded Waiting

- It must not be possible for a thread requiring access to a critical section to be delayed indefinitely.
- When no thread is in a critical section, any thread that requests entry must be permitted to enter without delay.

# Initial Attempts to Solve Problem

- Only 2 threads,  $T_0$  and  $T_1$
- General structure of thread  $T_i$  (other thread  $T_j$ )

do {

*enter section*

critical section

*exit section*

remainder section

} while (1);

- Threads may share some common variables to synchronize their actions.

# First Attempt: Algorithm 1

- Shared variables - initialization

```
int turn = 0;
```

- $\text{turn} == i \Rightarrow T_i$  can enter its critical section

- Thread  $T_i$

```
do{  
    while (turn != i) ;  
    critical section  
    turn = j;  
    reminder section  
} while(1);
```

- Satisfies mutual exclusion, but not progress

# Second Attempt: Algorithm 2

- Shared variables - initialization

```
int flag[2]; flag[0] = flag[1] = 0;
```

- $\text{flag}[i] == 1 \Rightarrow T_i$  can enter its critical section

- Thread  $T_i$

```
do{  
    flag[i] = 1;  
    while (flag[j] == 1) ;  
        critical section  
    flag[i] = 0;  
    remainder section  
} while(1);
```

- Satisfies mutual exclusion, but not progress requirement.

# Third Attempt: Algorithm 3

## (Peterson's Algorithm - 1981)

- Shared variables of algorithms 1 and 2 - initialization:

```
int flag[2]; flag[0] = flag[1] = 0;  
int turn = 0;
```

- Thread  $T_i$

```
do{  
  
    flag[i] = 1;  
    turn = j;  
    while ((flag[j] == 1) && turn == j) ;  
        critical section  
    flag[i] = 0;  
        remainder section  
} while(1);
```

- Solves the critical-section problem for two threads.



# Bakery Algorithm

## (Lamport 1979)

### A Solution to the Critical Section problem for $n$ threads

- Before entering its critical section, a thread receives a number. Holder of the smallest number enters the critical section.
- If threads  $T_i$  and  $T_j$  receive the same number, if  $i < j$ , then  $T_i$  is served first; else  $T_j$  is served first.
- The numbering scheme generates numbers in monotonically non-decreasing order; i.e., 1,1,1,2,3,3,3,4,4,5...

# Bakery Algorithm

- Notation “<” establishes lexicographical order among 2-tuples (ticket #, thread id #)

$(a,b) < (c,d)$  if  $a < c$  or if  $a == c$  and  $b < d$

$\max(a_0, \dots, a_{n-1}) = \{ k \mid k \geq a_i \text{ for } i = 0, \dots, n-1 \}$

- Shared data

```
int choosing[n];
```

```
int number[n];      - the ticket
```

Data structures are initialized to 0

# Bakery Algorithm

```
do{
    choosing[i] = 1;
    number[i] = max(number[0], number[1], ..., number[n-1]) + 1;
    choosing[i] = 0;
    for (j = 0; j < n; j++)
    {
        while (choosing[j] == 1) ;
        while ((number[j] != 0) &&
                ((number[j], j) < (number[i], i))) ;
    }
        critical section
    number[i] = 0;
        remainder section
} while(1);
```

# Mutual Exclusion - Hardware Support

## ● Interrupt Disabling

- Concurrent threads cannot overlap on a uniprocessor
- Thread will run until performing a system call or interrupt happens

## ● Special Atomic Machine Instructions

- Test and Set Instruction - read & write a memory location
- Exchange Instruction - swap register and memory location

## ● Problems with Machine-Instruction Approach

- Busy waiting
- Starvation is possible
- Deadlock is possible

# Synchronization Hardware

## • Test and modify the content of a word atomically

```
boolean TestAndSet(boolean &target) {  
    boolean rv = target;  
    target = true;  
    return rv;  
}
```

# Mutual Exclusion with Test-and-Set

- Shared data:

```
boolean lock = false;
```

- Thread  $T_i$

```
do{  
    while (TestAndSet(lock)) ;  
        critical section  
    lock = false;  
        remainder section  
} while(1);
```

# Synchronization Hardware

- Atomically swap two variables.

```
void Swap(boolean &a, boolean &b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```

# Mutual Exclusion with Swap

- Shared data (initialized to 0):

```
int lock = 0;
```

- Thread  $T_i$

```
int key;
```

```
do{
```

```
    key = 1;
```

```
    while (key == 1) Swap(lock, key);
```

```
        critical section
```

```
    lock = 0;
```

```
        remainder section
```

```
} while(1);
```



# Semaphores

- Semaphore  $S$  – integer variable
- can only be accessed via two atomic operations

*wait* ( $S$ ) :

    while ( $S \leq 0$ ) ;  
     $S--$ ;

*signal* ( $S$ ) :

$S++$ ;

# Critical Section of n Threads

## Shared data:

```
semaphore mutex; //initially mutex = 1
```

## Thread $T_i$ :

```
do{  
    wait(mutex);  
        critical section  
    signal(mutex);  
        remainder section  
} while(1);
```

# Semaphore Implementation

- Semaphores may suspend/resume threads
  - Avoid busy waiting
- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct thread *L;  
} semaphore;
```

- Assume two simple operations:
  - **suspend()** suspends the thread that invokes it.
  - **resume(*T*)** resumes the execution of a blocked thread *T*.

# Implementation

- Semaphore operations now defined as

*wait(S):*

```
S.value--;  
if (S.value < 0) {  
    add this thread to S.L;  
    suspend();  
}
```

*signal(S):*

```
S.value++;  
if (S.value <= 0) {  
    remove a thread T from S.L;  
    resume(T);  
}
```

# Two Types of Semaphores

- *Counting* semaphore
  - integer value can range over an unrestricted domain.
- *Binary* semaphore
  - integer value can range only between 0 and 1;
  - can be simpler to implement.
- Counting semaphore  $S$  can be implemented as a binary semaphore.

# Semaphore as a General Synchronization Tool

- Execute  $B$  in  $T_j$  only after  $A$  executed in  $T_i$
- Use semaphore  $flag$  initialized to 0
- Code:

$T_i$	$T_j$
...	...
$A$	$wait(flag)$
$signal(flag)$	$B$

# Deadlock and Starvation

- **Deadlock** – two or more threads are waiting indefinitely for an event that can be caused by only one of the waiting threads.
- Let S and Q be two semaphores initialized to 1

$T_0$	$T_1$
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
...	...
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q)</i>	<i>signal(S);</i>

- **Starvation** – indefinite blocking. A thread may never be removed from the semaphore queue in which it is suspended.
- Solution - all code should acquire/release semaphores in same order

# Windows Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems.
- Uses *spinlocks* on multiprocessor systems.
- Provides *dispatcher objects* which may act as mutexes and semaphores.
- Dispatcher objects may also provide *events*. An event acts much like a condition variable.



# Linux Synchronization

- Kernel *disables interrupts* for synchronizing access to global data on uniprocessor systems.
- Uses *spinlocks* for multiprocessor synchronization.
- Uses *semaphores* and *readers-writers* locks when longer sections of code need access to data.
- Implements POSIX synchronization primitives to support multitasking, multithreading (including real-time threads), and multiprocessing.

# Further Reading

- Ben-Ari, M., Principles of Concurrent Programming, Prentice Hall, 1982
- Lamport, L., The Mutual Exclusion Problem, Journal of the ACM, April 1986
- Abraham Silberschatz, Peter B. Galvin, and Greg Gagne, “*Operating System Concepts*”, John Wiley & Sons, 9th Ed., 2013
  - Chapter 5 - Process Synchronization
  - Chapter 7 - Deadlocks