

# Unit 4: Scheduling and Dispatch


## 4.2. Windows Processes and Threads

# Roadmap for Section 4.2.

- Windows Processes and Threads
- Performance Counters
- Process and Thread Lifetime
- Jobs
- Windows APIs for Process and Thread creation

# Windows Processes

- What is a process?
  - Represents an instance of a running program
    - you create a process to run a program
    - starting an application creates a process
  - Process defined by:
    - Address space
    - Resources (e.g. open handles)
    - Security profile (token)
- Every process starts with one thread
  - First thread executes the program's "main" function
    - Can create other threads in the same process
    - Can create additional processes



# Windows Threads

## What is a thread?

- An execution context within a process
- Unit of scheduling (threads run, processes don't run)
- All threads in a process share the same per-process address space
  - Services provided so that threads can synchronize access to shared resources (critical sections, mutexes, events, semaphores)
- All threads in the system are scheduled as peers to all others, without regard to their "parent" process

## System calls

- Primary argument to `CreateProcess()` is image file name (or command line)
- Primary argument to `CreateThread()` is a function entry point address

# Processes & Threads

- Why divide an application into multiple threads?
  - Perceived user responsiveness, parallel/background execution
    - Examples: Word background print – can continue to edit during print
  - Take advantage of multiple processors
    - On an MP system with  $n$  CPUs,  $n$  threads can literally run at the same time
    - Question: given a single threaded application, will adding a 2nd processor make it run faster?
  - Does add complexity
    - Synchronization
    - Scalability well is a different question...
      - # of multiple runnable threads vs # CPUs
      - Having too many runnable threads causes excess context switching

# Per-Process Data

- Each process has its own...
  - Virtual address space (including program code, global storage, heap storage, threads' stacks)
    - processes cannot corrupt each other's address space by mistake
  - Working set (physical memory “owned” by the process)
  - Access token (includes security identifiers)
  - Handle table for Windows kernel objects
  - Environment strings
  - Command line
  - These are common to all threads in the process, but separate and protected between processes

# Per-Thread Data

- Each thread has its own...
  - User-mode stack (arguments passed to thread, automatic storage, call frames, etc.)
  - Kernel-mode stack (for system calls)
  - Thread Local Storage (TLS) – array of pointers to allocate unique data
  - Scheduling state (Wait, Ready, Running, etc.) and priority
  - Hardware context (saved in CONTEXT structure if not running)
    - Program counter, stack pointer, register values
    - Current access mode (user mode or kernel mode)
  - Access token (optional – overrides process's implicit access token if present)

# Process and Thread Identifiers

- Every process and every thread has an identifier
- Generically: “client ID” (debugger shows as “CID”)
  - A.K.A. “process ID” and “thread ID”, respectively
  - Process IDs and thread IDs are in the same “number space”
  - These identify the requesting process or thread to its subsystem “server” process (i.e. CSRSS.EXE), in API calls that need the server’s help
- Visible in PerfMon, Task Manager (for processes), Process Viewer (for processes), kernel debugger, etc.
- IDs are unique among all existing processes and threads
  - But might be reused as soon as a process or thread is deleted



# Process-Related Performance Counters

Object: Counter	Function
Process:%PrivilegedTime	Percentage of CPU time that the threads in the process have run in kernel mode
Process:%ProcessorTime	Percentage of CPU time that threads have used during specified interval, i.e. %PrivilegedTime + %UserTime
Process:%UserTime	Percentage of CPU time that the threads in the process have run in user mode
Process: ElapsedTime	Total lifetime of process in seconds
Process: ID Process	PID – process ID
Process: ThreadCount	Number of threads in a process

# Thread-Related Performance Counters

Object: Counter	Function
Thread:%PrivilegedTime	Percentage of CPU time that the thread was run in kernel mode
Thread:%ProcessorTime	Percentage of CPU time that the thread has used during specified interval, i.e. $\%PrivilegedTime + \%UserTime$
Thread:%UserTime	Percentage of CPU time that the thread has run in user mode
Thread: ElapsedTime	Total lifetime of process in seconds
Thread: ID Process	PID – thread's process ID
Thread: ID Thread	TID – thread ID

# Thread-Related Performance Counters (contd.)

Object: Counter	Function
<b>Process:</b> Priority Base	Base priority of process: starting priority for all threads within that process
Thread: Priority Base	Base priority of thread: may differ from the thread's starting priority
Thread: Priority Current	The thread's current dynamic priority
Thread: Start Address	The thread's starting virtual address (the same for most threads)
Thread: Thread State	Value from 0 through 7 – current state of thread
Thread: Thread Wait Reason	Value from 0 through 19 – reason why the thread is in wait state

# Tools for Obtaining Process & Thread Information

- Many overlapping tools (most show one item the others do not)
- Built-in tools introduced in Windows 2000/XP:
  - **Task Manager**, **Performance Tool**
  - Tasklist (new in XP)
- Support Tools
  - pviewer - process and thread details (GUI)
  - pmon - process list (character cell)
  - tlist - shows process tree and thread details (character cell)
- Resource Kit tools:
  - apimon - system call and page fault monitoring (GUI)
  - oh - display open handles (character cell)
  - pviewer - processes and threads and security details (GUI)
  - ptree - display process tree and kill remote processes (GUI)
  - pulist - lists processes and usernames (character cell)
  - pstat - process/threads and driver addresses (character cell)
  - qslice - can show process-relative thread activity (GUI)
- Tools from [www.sysinternals.com](http://www.sysinternals.com)
  - **Process Explorer** – super Task Manager – shows open files, loaded DLLs, security info, etc.
  - Pslist – list processes on local or remote systems
  - Ntpmon – shows process/thread create/deletes (and context switches on MP systems only)
  - Listdlls – displays full path of EXE & DLLs loaded in each process

# Process Lifetime

- Created as an empty shell
- Address space created with only ntdll and the main image unless created by POSIX fork()
- Handle table created empty or populated via duplication from parent
- Process is partially destroyed on last thread exit
- Process totally destroyed on last dereference

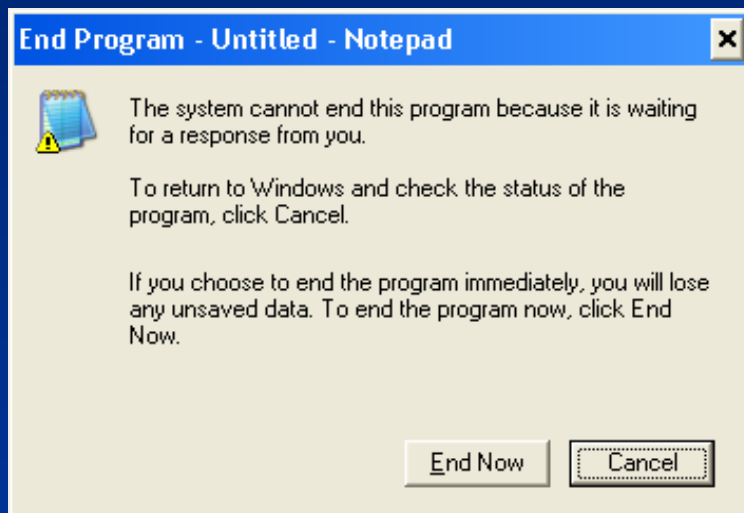
# Thread Lifetime

- Created within a process with a CONTEXT record
  - Starts running in the kernel but has a trap frame to return to user mode
- Threads run until they:
  - The thread returns to the OS
  - ExitThread is called by the thread
  - TerminateThread is called on the thread
  - ExitProcess is called on the process

# Why/When Do Processes Exit? (or Terminate?)

- Normal: Application decides to exit (ExitProcess)

- Usually due to a request from the UI
- or: C RTL does ExitProcess when primary thread function (main, WinMain, etc.) returns to caller
  - this forces TerminateThread on the process's remaining threads
  - or, any thread in the process can do an explicit ExitProcess



- Orderly exit requested from the desktop (ExitProcess)

- e.g. "End Task" from Task Manager "Tasks" tab
- Task Manager sends a WM\_CLOSE message to the window's message loop...
- ...which should do an ExitProcess (or equivalent) on itself

- Forced termination (TerminateProcess)


- if no response to "End Task" in five seconds, Task Manager presents End Program dialog (which does a TerminateProcess)
- or: "End Process" from Task Manager Processes tab

- Unhandled exception

- Covered in Unit 4.3 (Process and Thread Internals)



# Jobs



- Jobs are collections of processes
  - Can be used to specify limits on CPU, memory, and security
  - Enables control over some unique process & thread settings not available through any process or thread system call
    - E.g. length of thread's timeslice (i.e. quantum length)
- How do processes become part of a job?
  - Job object has to be created (CreateJobObject)
  - Then processes are explicitly added (AssignProcessToJob)
    - Processes created by processes in a job automatically are part of the job
      - Unless restricted, processes can “break away” from a job
  - Then quotas and limits are defined (SetInformationJobObject)
    - Examples on next slide...



# Job Settings

- Quotas and restrictions:

- Quotas: total CPU time, # active processes, per-process CPU time, memory usage
- Run-time restrictions: priority of all the processes in job, processors on which threads in job can run, etc.
- Security restrictions: limits what processes can do, e.g.
  - Not acquire administrative privileges
  - Not accessing windows outside the job, no reading/writing the clipboard
- Scheduling class: number from 0-9 (5 is default) – this affects the length of thread's timeslice (or quantum)
  - E.g. can be used to achieve “class scheduling” (partition CPU)

# Jobs

- Examples where Windows OS uses jobs:
  - Add/Remove Programs (“ARP Job”)
  - WMI provider
  - RUNAS service (SecLogon) uses jobs to terminate processes at log out
- Process Explorer highlights processes that are members of jobs
  - Color can be configured with Options->Configure Highlighting
  - For processes in a job, click on Job tab in process properties to see details



# Process Windows APIs

- CreateProcess
- OpenProcess
- GetCurrentProcessId – returns a global ID (“client” ID)
- GetCurrentProcess – returns a handle
- ExitProcess
- TerminateProcess – no DLL notification
- Get / SetProcessShutdownParameters
- GetExitCodeProcess
- GetProcessTimes
- GetStartupInfo



# Windows Thread APIs

- CreateThread
- CreateRemoteThread
- GetCurrentThreadId – returns global ID (“client” ID)
- GetCurrentThread – returns handle
- SuspendThread / ResumeThread
- ExitThread
- TerminateThread – no DLL notification
- GetExitCodeThread
- GetThreadTimes
- Windows 2000 added:
  - OpenThread
  - new thread pooling APIs

# Fibers

- Implemented completely in user mode
  - no “internals” ramifications
  - Fibers are still scheduled as threads
  - Fiber APIs allow different execution contexts within a (kernel) thread
    - stack
    - fiber-local storage
    - some registers (essentially those saved and restored for a procedure call)
    - cooperatively “scheduled” within the thread
  - Analogous to threading libraries under many Unix systems
  - Analogous to co-routines in assembly language
  - Allow easy porting of apps that “did their own threads” under other systems

# Process Creation

- No parent/child relation in Win32 environment subsystem
- *CreateProcess()* – new process with primary thread

```
BOOL CreateProcess(  
    LPCSTR lpApplicationName,  
    LPSTR lpCommandLine,  
    LPSECURITY_ATTRIBUTES lpProcessAttributes,  
    LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    BOOL bInheritHandles,  
    DWORD dwCreationFlags,  
    LPVOID lpEnvironment,  
    LPCSTR lpCurrentDirectory,  
    LPSTARTUPINFO lpStartupInfo,  
    LPPROCESS_INFORMATION lpProcessInformation)
```

# Parameters

- dwCreationFlags:
  - CREATE\_SUSPENDED, DETACHED\_PROCESS, CREATE\_NEW\_CONSOLE, CREATE\_NEW\_PROCESS\_GROUP
- lpStartupInfo:
  - Main window appearance
  - Parent's info: GetStartupInfo
  - hStdIn, hStdOut, hStdErr fields for I/O redirection
- lpProcessInformation:
  - Pointer to handle & ID of new process/thread

```
typedef struct _PROCESS_INFORMATION {  
    HANDLE hProcess;  
    HANDLE hThread;  
    DWORD dwProcessId;  
    DWORD dwThreadId;  
} PROCESS_INFORMATION;
```

# UNIX & Win32 comparison

- Windows API has no equivalent to `fork()`
- `CreateProcess()` similar to `fork()` + `exec()` in child
- UNIX `$PATH` vs. `lpCommandLine` argument
  - Win32 searches in directory of current process image; in current directory; in Windows system directory (`GetSystemDirectory`); in Windows directory (`GetWindowsDirectory`); in directories given in `PATH`
- Windows API has no parent/child relations for processes
- No UNIX process groups in Windows API
  - Limited form: group = processes to receive a console event



# Windows API Thread Creation

```
HANDLE CreateThread (  
    LPSECURITY_ATTRIBUTES lpsa,  
    DWORD cbStack,  
    LPTHREAD_START_ROUTINE lpStartAddr,  
    LPVOID lpvThreadParm,  
    DWORD fdwCreate,  
    LPDWORD lpIDThread)
```

cbStack == 0: thread's  
stack size defaults to  
primary thread's size

- lpstartAddr points to function declared as  
`DWORD WINAPI ThreadFunc(LPVOID)`
- lpvThreadParm is 32-bit argument
- lpIDThread points to DWORD that receives thread ID  
non-NULL pointer !

# Exiting and Terminating a Process

- Shared resources must be freed before exiting
  - Mutexes, semaphores, events
  - Use structured exception handling

• But:

`_finally`, `_except`  
handlers are not  
executed on  
`ExitProcess`;

- no SEH on  
`TerminateProcess`

```
VOID ExitProcess(  
    UINT uExitCode);  
  
BOOL TerminateProcess(  
    HANDLE hProcess,  
    UINT uExitCode);  
  
BOOL GetExitCodeProcess(  
    HANDLE hProcess,  
    LPDWORD lpExitCode);
```

# Windows API Thread Termination

```
VOID ExitThread( DWORD devExitCode )
```

- When the last thread in a process terminates, the process itself terminates  
(TerminateThread() does not execute final SEH)
- Thread continues to exist until last handle is closed  
(CloseHandle())

```
BOOL GetExitCodeThread (  
    HANDLE hThread, LPDWORD lpdwExitCode)
```

- Returns exit code or STILL\_ACTIVE

# Suspending and Resuming Threads

- Each thread has suspend count
- Can only execute if suspend count == 0
- Thread can be created in suspended state

```
DWORD ResumeThread (HANDLE hThread)  
DWORD SuspendThread(HANDLE hThread)
```

- Both functions return suspend count or 0xFFFFFFFF on failure

# Synchronization & Remote Threads

- WaitForSingleObject() and WaitForMultipleObjects() with thread handles as arguments perform thread synchronization
  - Waits for thread to become signaled
  - ExitThread(), TerminateThread(), ExitProcess() set thread objects to signaled state
- CreateRemoteThread() allows creation of thread in another process
  - Not implemented in Windows 9x
- C library is not thread-safe; use libcmt.lib instead
  - #define \_MT before any include
  - Use \_beginthreadex/\_endthreadex instead of Create/ExitThread

# Further Reading

- Pavel Yosifovich, Alex Ionescu, et al., “Windows Internals”, 7th Edition, Microsoft Press, 2017.
  - Chapter 3 – Processes and jobs (from pp. 156)
  - Chapter 4 – Threads (from pp. 275)
- Jeffrey Richter, Programming Applications for Microsoft Windows, 4th Edition, Microsoft Press, September 1999.
  - Chapter 4 - Processes
  - Chapter 5 - Jobs
  - Chapter 6 - Thread Basics