

Unit 4: Scheduling and Dispatch

4.4. Windows Thread Scheduling

Roadmap for Section 4.4.

- Windows Scheduling Principles
- Windows API vs. NT Kernel Priorities
- Scheduling Data Structures
- Scheduling Scenarios
- Priority Boosts and Priority Adjustments

Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – # of processes/threads that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process/thread
- Waiting time – amount of time a process/thread has been waiting in the ready queue
- Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (i.e., *the hourglass*)

How does the Windows scheduler relate to the issues discussed:

- Priority-driven, preemptive scheduling system
- Highest-priority runnable thread always runs
- Thread runs for time amount of *quantum*
- No single scheduler – event-based scheduling code spread across the kernel
- Dispatcher routines triggered by the following events:
 - Thread becomes ready for execution
 - Thread leaves running state (quantum expires, wait state)
 - Thread's priority changes (system call/NT activity)
 - Processor affinity of a running thread changes

Windows Scheduling Principles

- 32 priority levels
- Threads within same priority are scheduled following the Round-Robin policy
- Non-Realtime Priorities are adjusted dynamically
 - Priority elevation as response to certain I/O and dispatch events
 - Quantum stretching to optimize responsiveness
- Realtime priorities (i.e.; > 15) are assigned statically to threads

Scheduling

- Multiple threads may be ready to run
- “Who gets to use the CPU?”
- From Windows API point of view:
 - Processes are given a priority class upon creation
 - Idle, Normal, High, Realtime
 - Windows 2000 added “Above normal” and “Below normal”
 - Threads have a relative priority within the class
 - Idle, Lowest, Below_Normal, Normal, Above_Normal, Highest, and Time_Critical
- From the kernel’s view:
 - Threads have priorities 0 through 31
 - Threads are scheduled, not processes
 - Process priority class is not used to make scheduling decisions

Windows Scheduling-related APIs:

Get/SetPriorityClass

Get/SetThreadPriority

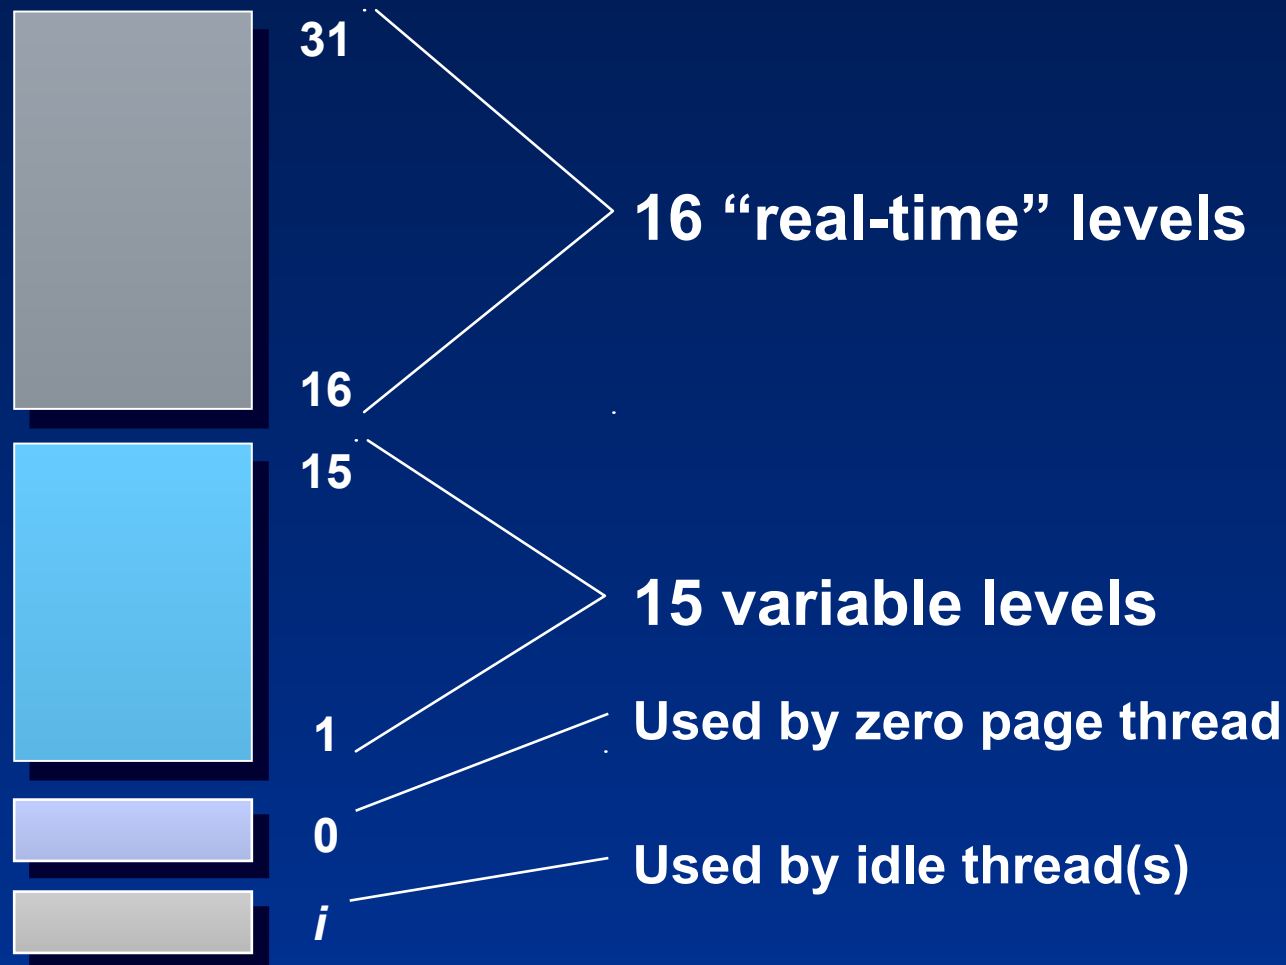
Get/SetProcessAffinityMask

SetThreadAffinityMask

SetThreadIdealProcessor

Suspend/ResumeThread

Kernel: Thread Priority Levels



Windows vs. NT Kernel Priorities

		Win32 Process Classes					
		Realtime	High	Above Normal	Normal	Below Normal	Idle
Win32 Thread Priorities	Time-critical	31	15	15	15	15	15
	Highest	26	15	12	10	8	6
	Above-normal	25	14	11	9	7	5
	Normal	24	13	10	8	6	4
	Below-normal	23	12	9	7	5	3
	Lowest	22	11	8	6	4	2
	Idle	16	1	1	1	1	1

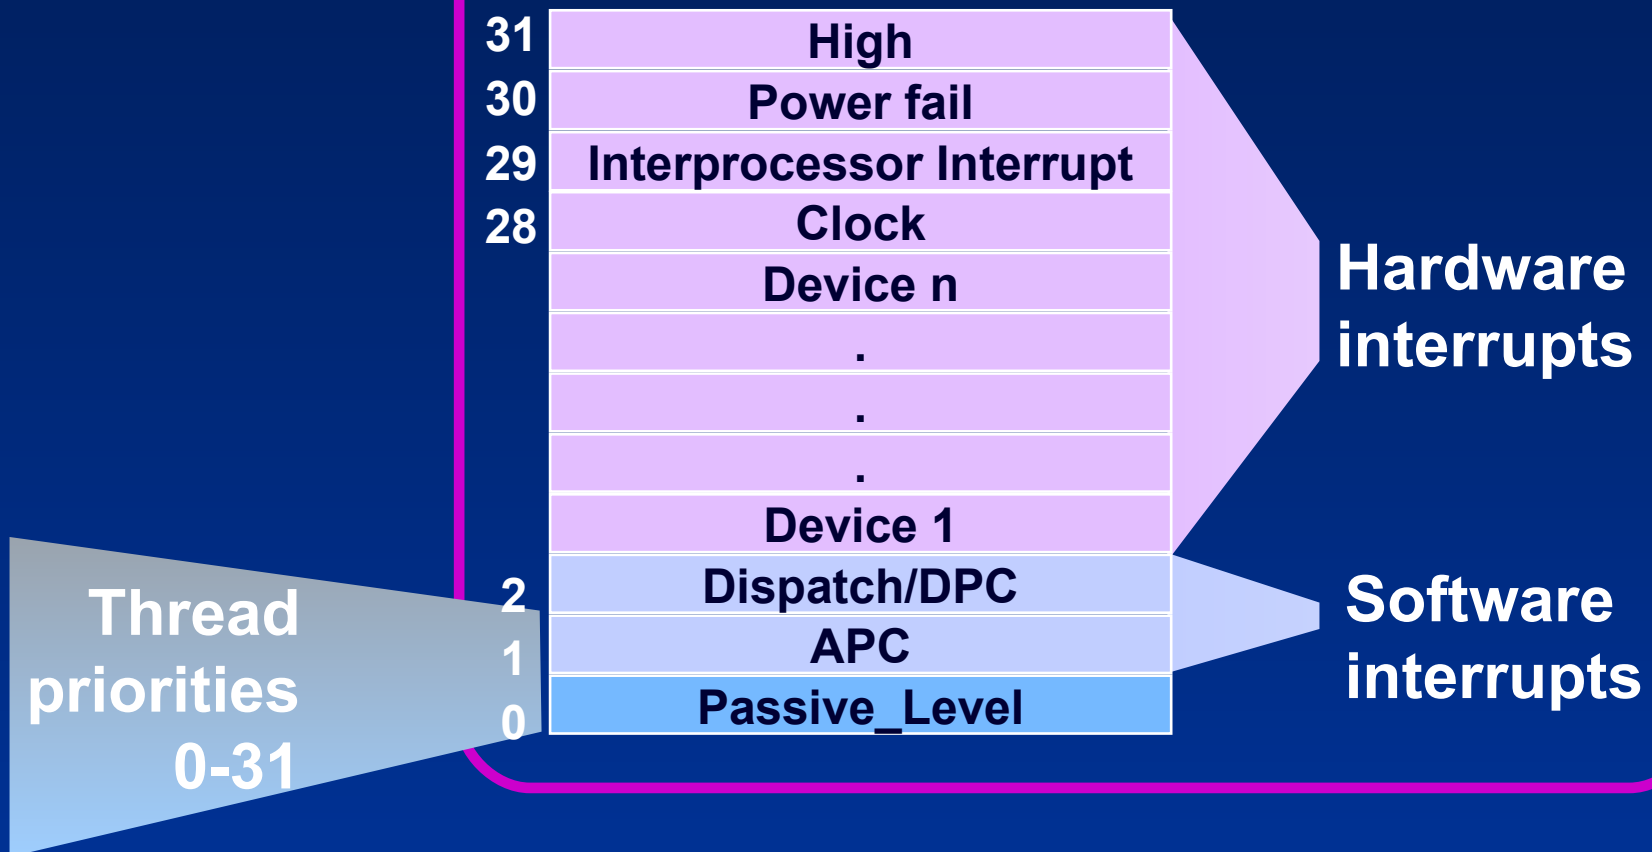
- Table shows base priorities (“current” or “dynamic” thread priority may be higher if base is < 15)
- Many utilities (such as Process Viewer) show the “dynamic priority” of threads rather than the base (Performance Monitor can show both)
- Drivers can set to any value with KeSetPriorityThread

Special Thread Priorities

- Idle threads -- one per CPU
 - When no threads want to run, Idle thread “runs”
 - Not a real priority level - appears to have priority zero, but actually runs “below” priority 0
 - Provides CPU idle time accounting (unused clock ticks are charged to the idle thread)
 - Loop:
 - Calls HAL to allow for power management
 - Processes DPC list
 - Dispatches to a thread if selected
 - Server 2003: in certain cases, scans per-CPU ready queues for next thread
- Zero page thread -- one per NT system
 - Zeroes pages of memory in anticipation of “demand zero” page faults
 - Runs at priority zero (lower than any reachable from Windows)
 - Part of the “System” process (not a complete process)

Thread Scheduling Priorities vs. Interrupt Request Levels (IRQs)

IRQLs (x86)



Single Processor Thread Scheduling

- Priority driven, preemptive
 - 32 queues (FIFO lists) of “ready” threads
 - UP: highest priority thread always runs
 - MP: One of the highest priority runnable thread will be running somewhere
 - No attempt to share processor(s) “fairly” among processes, only among threads
 - Time-sliced, round-robin within a priority level
- Event-driven; no guaranteed execution period before preemption
 - When a thread becomes Ready, it either runs immediately or is inserted at the tail of the Ready queue for its current (dynamic) priority

Thread Scheduling

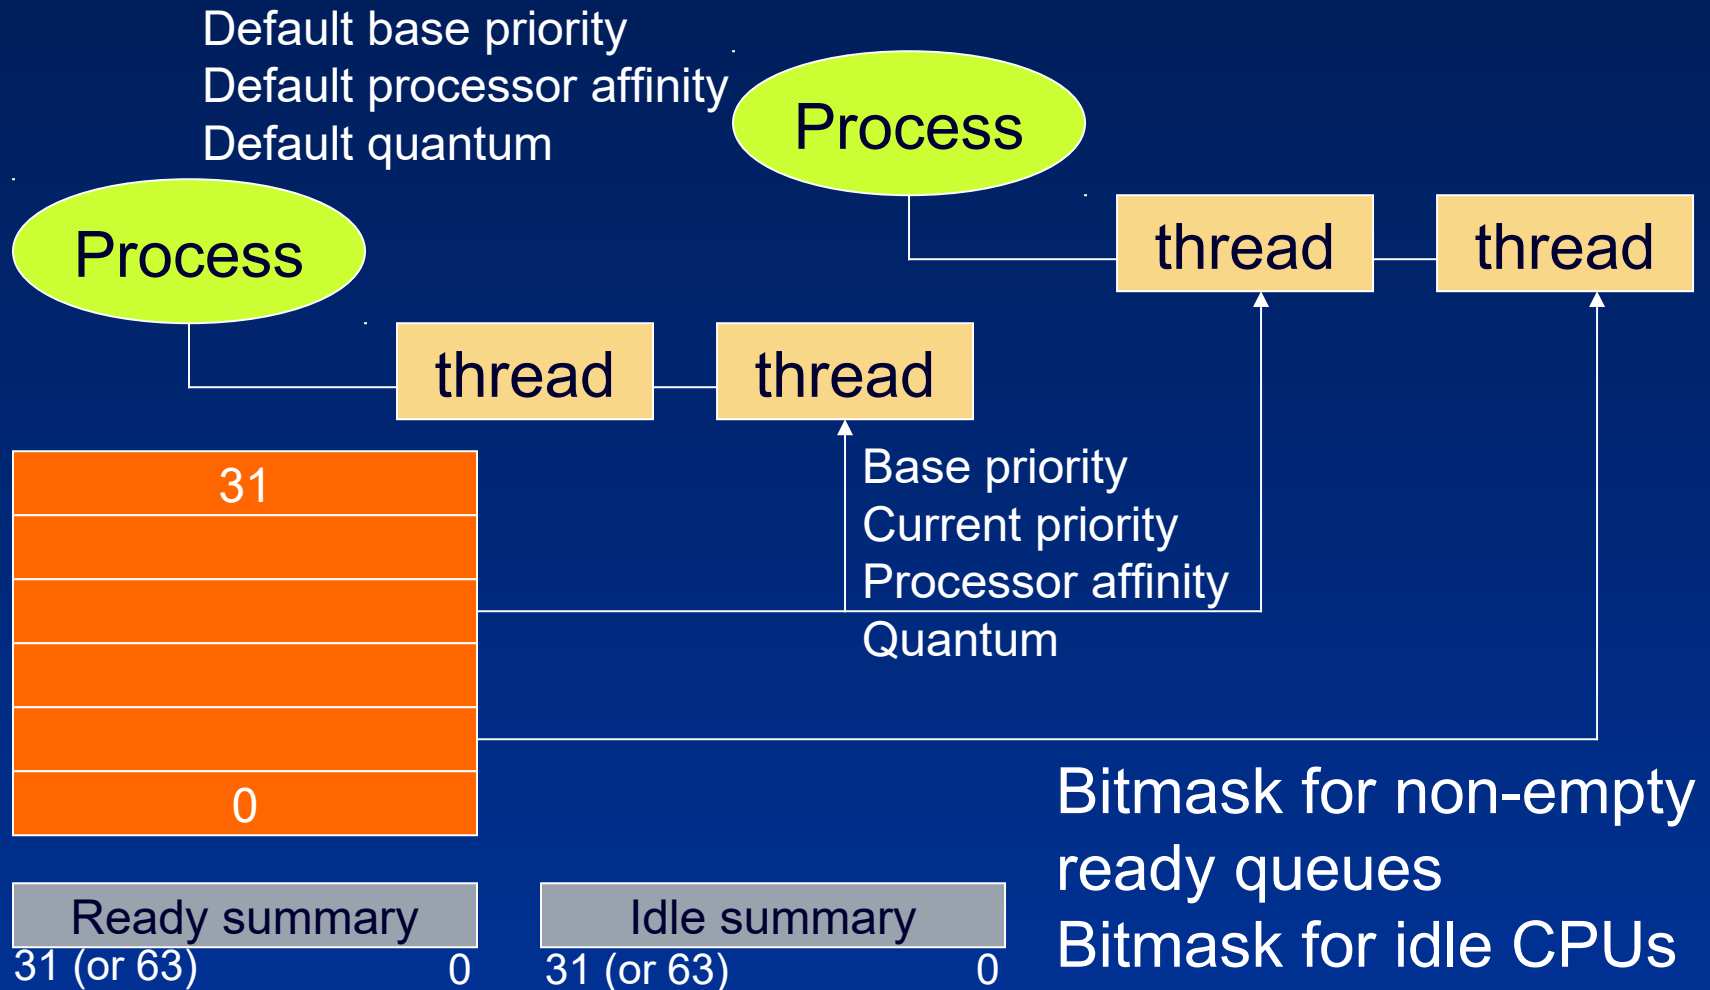
- No central scheduler!

- i.e. there is no always-instantiated routine called “the scheduler”
- The “code that does scheduling” is not a thread
- Scheduling routines are simply called whenever events occur that change the Ready state of a thread
- Things that cause scheduling events include:
 - interval timer interrupts (for quantum end)
 - interval timer interrupts (for timed wait completion)
 - other hardware interrupts (for I/O wait completion)
 - one thread changes the state of a waitable object upon which other thread(s) are waiting
 - a thread waits on one or more dispatcher objects
 - a thread priority is changed

- Based on doubly-linked lists (queues) of Ready threads

- Nothing that takes “order- n time” for n threads

Scheduling Data Structures



Scheduling Scenarios

• Preemption

- A thread becomes Ready at a higher priority than the running thread
- Lower-priority Running thread is preempted
- Preempted thread goes back to head of its Ready queue
 - action: pick lowest priority thread to preempt

• Voluntary switch

- Waiting on a dispatcher object
- Termination
- Explicit lowering of priority
 - action: scan for next Ready thread (starting at your priority & down)


• Running thread experiences quantum end

- Priority is decremented unless already at thread base priority
- Thread goes to tail of ready queue for its new priority
- May continue running if no equal or higher-priority threads are Ready
 - action: pick next thread at same priority level

Scheduling Scenarios

Preemption

- Preemption is strictly event-driven
 - does not wait for the next clock tick
 - no guaranteed execution period before preemption
 - threads in kernel mode may be preempted (unless they raise IRQL to ≥ 2)




- A preempted thread goes back to the head of its ready queue

Scheduling Scenarios

Ready after Wait Resolution


- If newly-ready thread is not of higher priority than the running thread...
- ...it is put at the tail of the ready queue for its current priority
 - If priority ≥ 14 quantum is reset (t.b.d.)
 - If priority < 14 and you're about to be boosted and didn't already have a boost, quantum is set to process quantum - 1



Scheduling Scenarios

Voluntary Switch

- When the running thread gives up the CPU...
- ...Schedule the thread at the head of the next non-empty “ready” queue




to Waiting state

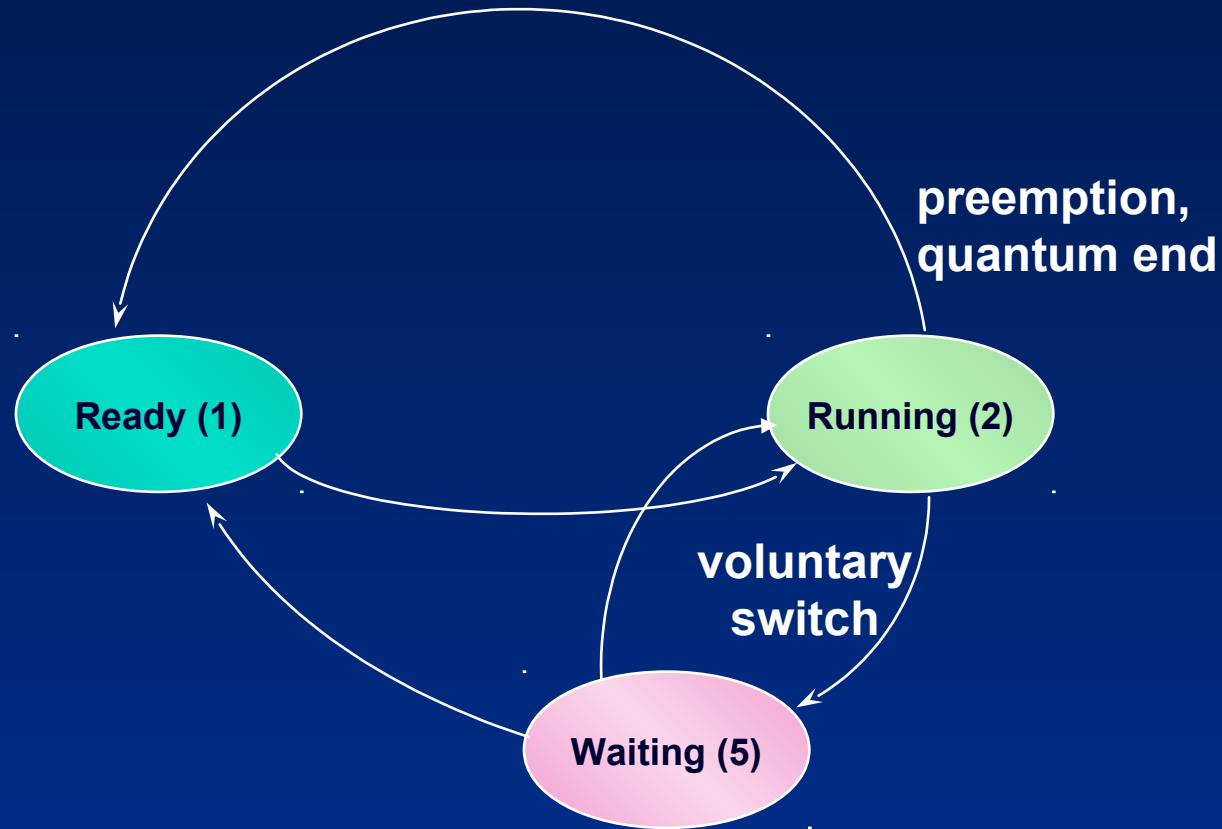
Scheduling Scenarios

Quantum End (“time-slicing”)

- When the running thread exhausts its CPU quantum, it goes to the end of its ready queue
 - Applies to both real-time and dynamic priority threads, user and kernel mode
 - Quantum can be disabled for a thread by a kernel function
 - Default quantum on NT Client versions is 2 clock ticks, 12 on Server versions
 - standard clock tick is 10 msec; might be 15 msec on some MP Pentium systems
 - if no other ready threads at that priority, same thread continues running (just gets new quantum)
 - if running at boosted priority, priority decays by one at quantum end (described later)



Basic Thread Scheduling States



Priority Adjustments

- Dynamic priority adjustments (boost and decay) are applied to threads in “dynamic” classes
 - Threads with base priorities 1-15 (technically, 1 through 14)
 - Disable if desired with `SetThreadPriorityBoost` or `SetProcessPriorityBoost`
- Five types:
 - I/O completion
 - Wait completion on events or semaphores
 - When threads in the foreground process complete a wait
 - When GUI threads wake up for windows input
 - For CPU starvation avoidance
- No automatic adjustments in “real-time” class (16 or above)
 - “Real time” here really means “system won’t change the relative priorities of your real-time threads”
 - Hence, scheduling is predictable with respect to other “real-time” threads (but not for absolute latency)

Priority Boosting

To favor I/O intense threads:

- After an I/O: specified by device driver
 - `IoCompleteRequest(Irp, PriorityBoost)`

Common boost values (see NTDDK.H)

1: disk, CD-ROM, parallel, video

2: serial, network, named pipe, mailslot

6: keyboard or mouse

8: sound

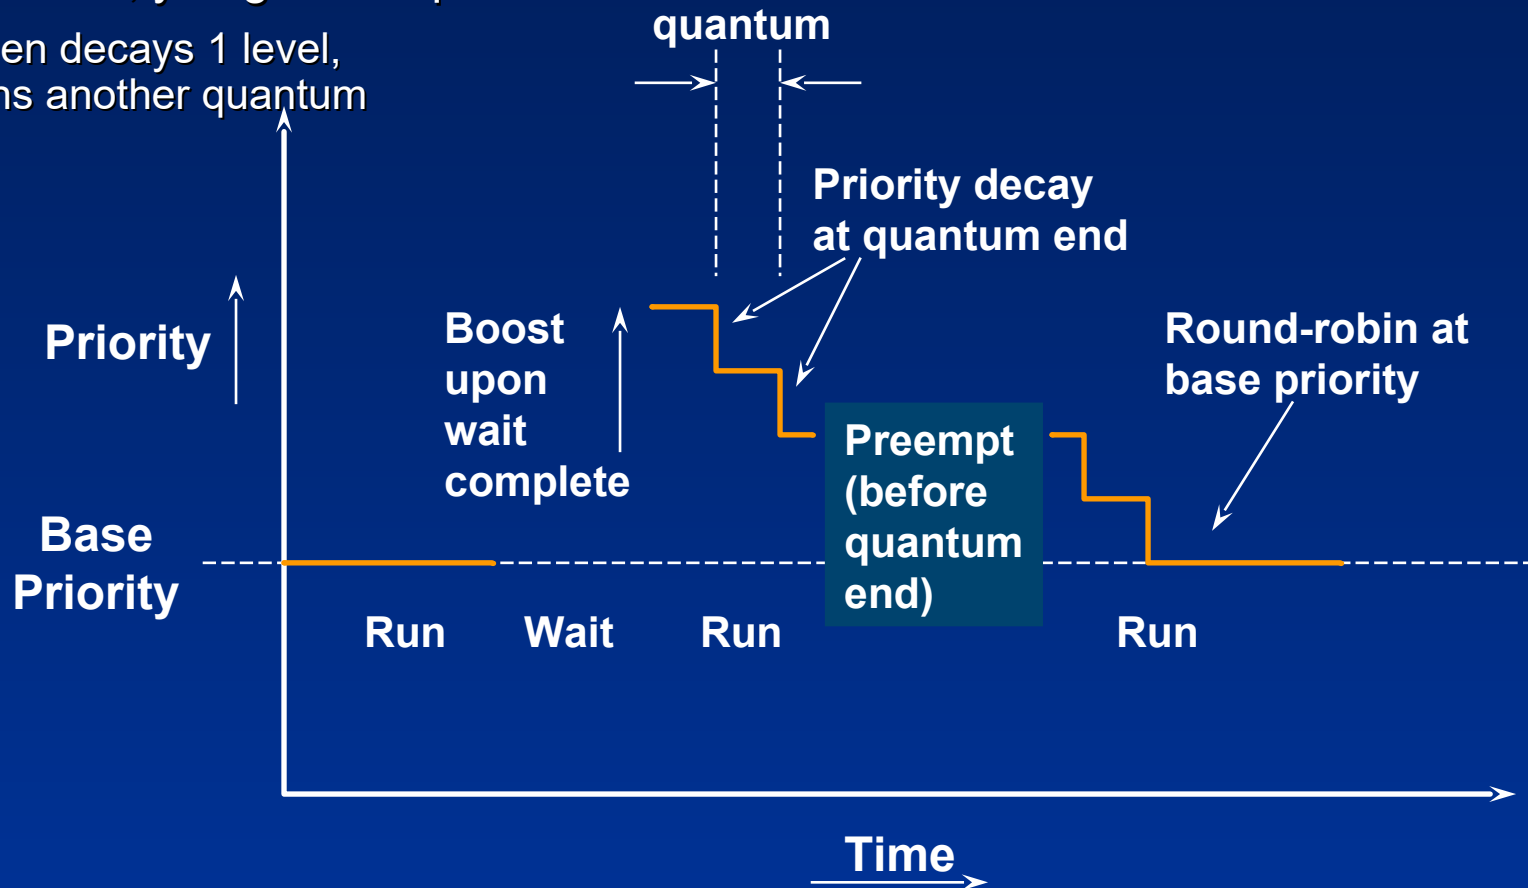
Other cases discussed in the Windows Scheduling Internals Section

- After a wait on executive event or semaphore
- After any wait on a dispatcher object by a thread in the foreground process
- GUI threads that wake up to process windowing input (e.g. windows messages) get a boost of 2

Thread Priority Boost and Decay

Behavior of these boosts:

- Applied to thread's base priority
 - will not take you above priority 15
- After a boost, you get one quantum
 - Then decays 1 level, runs another quantum



Further Reading

- Pavel Yosifovich, Alex Ionescu, et al., “Windows Internals”, 7th Edition, Microsoft Press, 2017.
 - Chapter 4 – Threads (from pp. 275)
 - Thread scheduling (from pp. 303)
 - Thread states (from pp. 315)
 - Scheduling scenarios (from pp. 361)