

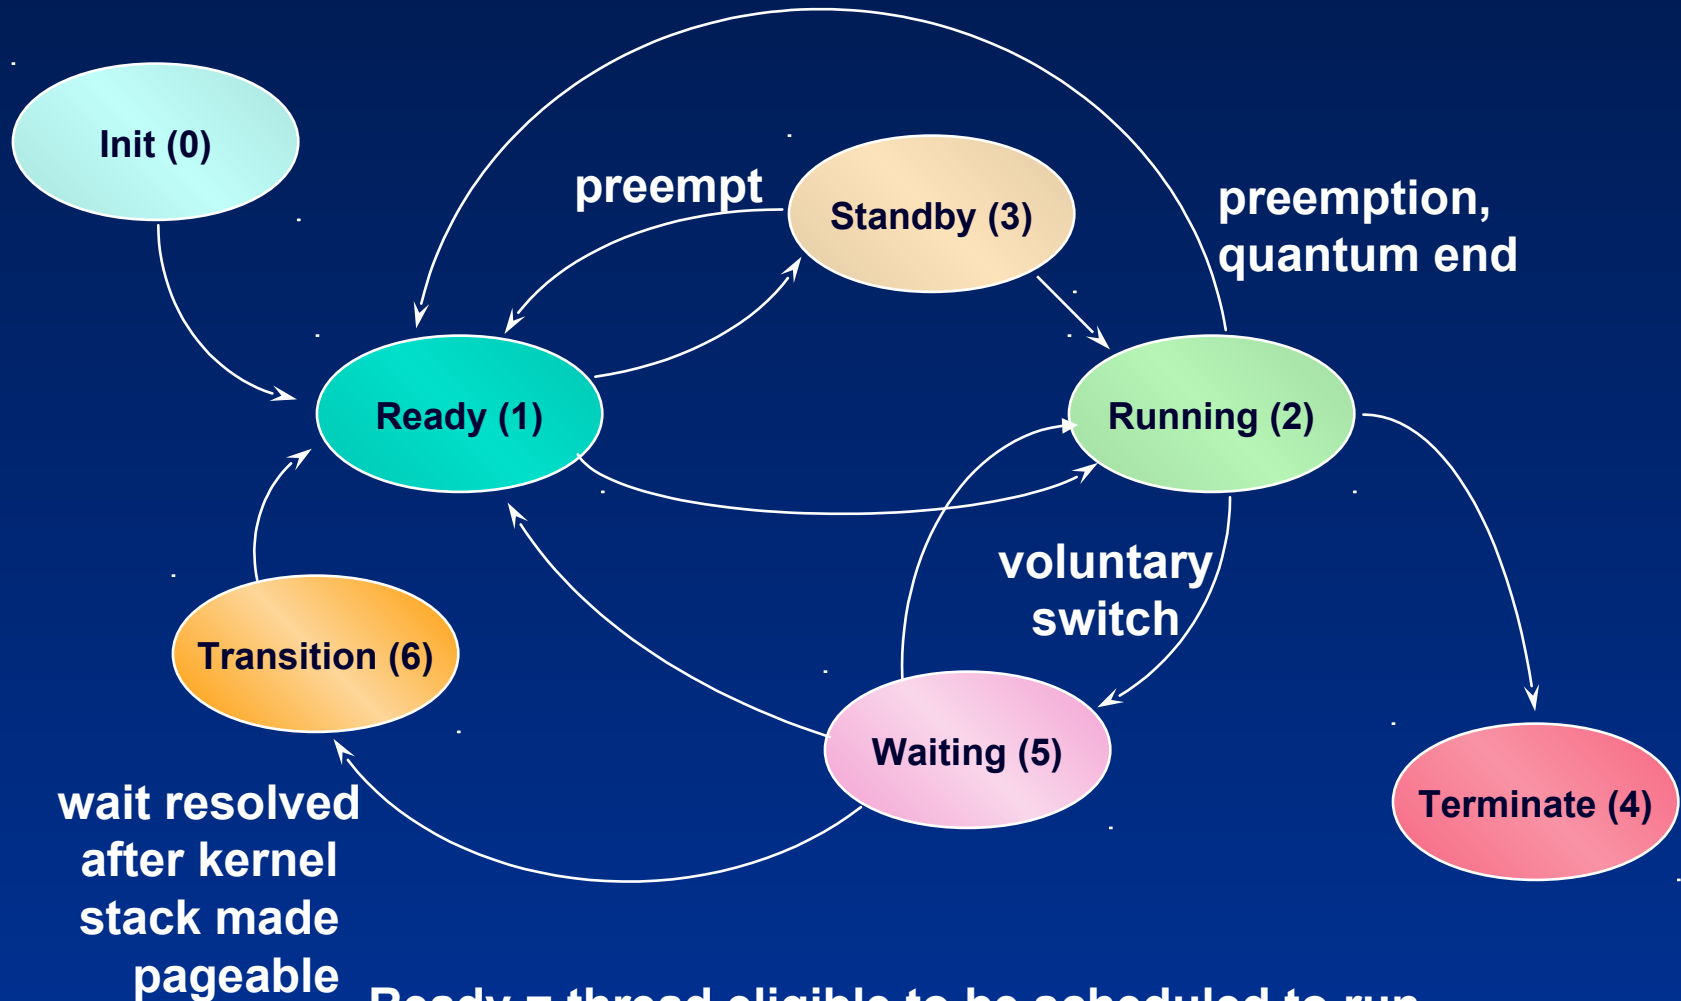
Unit 4: Scheduling and Dispatch

4.5. Advanced Windows Thread Scheduling

Roadmap for Section 4.5.

- Thread Scheduling Details
- Quantum Stretching
- CPU Starvation Avoidance
- Multiprocessor Scheduling
- Windows Server 2003 Enhancements

Thread Scheduling States (2000, XP)



Ready = thread eligible to be scheduled to run
Standby = thread is selected to run on another CPU

Other Thread States

Transition

- Thread was in a wait entered from user mode for 12 seconds or more
- System was short on physical memory
- Balance set manager (t.b.d.) marked the thread's kernel stack as pageable (preparatory to "outswapping" the thread's process)
- Later, the thread's wait was satisfied, but...
- ...Thread can't become Ready until the system allocates a nonpageable kernel stack; it is in the "transition" state until then

Initiate

- Thread is "under construction" and can't run yet

Standby

- One processor has selected a thread for execution on another processor

Terminate

- Thread has executed its last code, but can't be deleted until all handles and references to it are closed (Object Manager)

Scheduling Scenarios

Quantum Details

- Quantum internally stored as “3 * number of clock ticks”
 - Default quantum is 6 on NT Client versions, 36 on Server versions
- Thread->Quantum field is decremented by 3 on every clock tick
- Process and thread objects have a Quantum field
 - Process quantum is simply used to initialize thread quantum for all threads in the process
- Quantum decremented by 1 when you come out of a wait
 - So that threads that get boosted after I/O completion won't keep running and never experiencing quantum end
 - Prevents I/O bound threads from getting unfair preference over CPU bound threads

Scheduling Scenarios

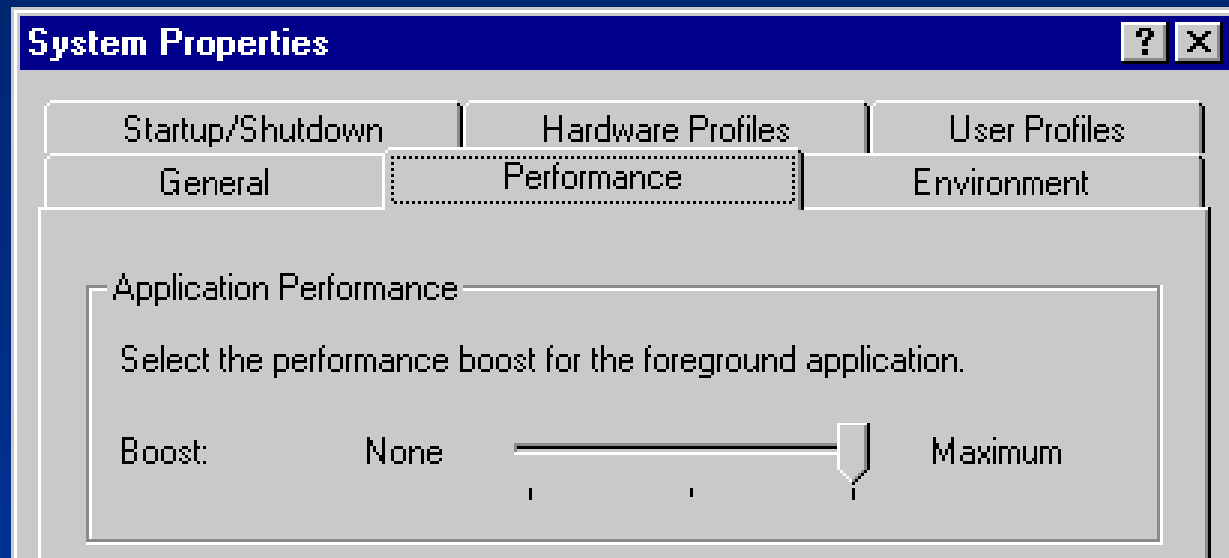
Quantum Details

- When Thread->Quantum reaches zero (or less than zero):
 - you've experienced quantum end
 - Thread->Quantum = Process->Quantum; // restore quantum
 - for dynamic-priority threads, this is the only thing that restores the quantum
 - for real-time threads, quantum is also restored upon preemption
- Interval timer interrupts when previous IRQL ≥ 2 :
 - are not charged to the current thread's "privileged" time
 - but do cause the thread "remaining quantum" counter to be decremented

Quantum Stretching

(favoring foreground applications)

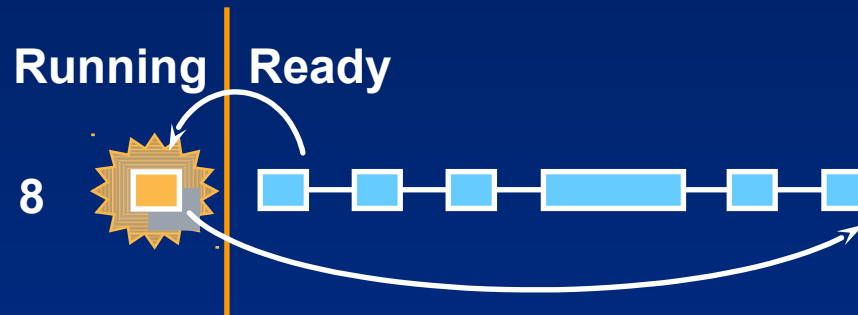
- If normal-priority process owns the foreground window, its threads may be given longer quantum
 - Set by Control Panel / System applet / Performance tab
 - Stored in ...\\System\\CurrentControlSet\\Control\\PriorityControl
Win32PrioritySeparation = 0, 1, or 2
 - New behavior with NT 4.0 (formerly implemented via priority shift)



Screen snapshot from:
Control Panel | System |
Performance tab

Quantum Stretching

- Resulting quantum:
 - “Maximum” = 6 ticks
 - (middle) = 4 ticks
 - “None” = 2 ticks



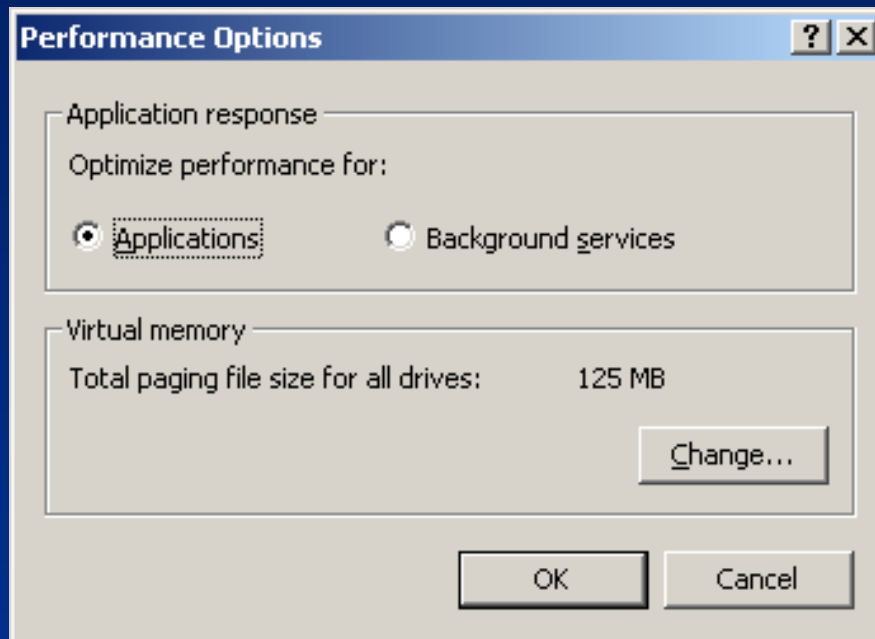
- Quantum stretching does not happen on NT Server versions
 - Quantum on Server is always 12 ticks

Quantum Selection



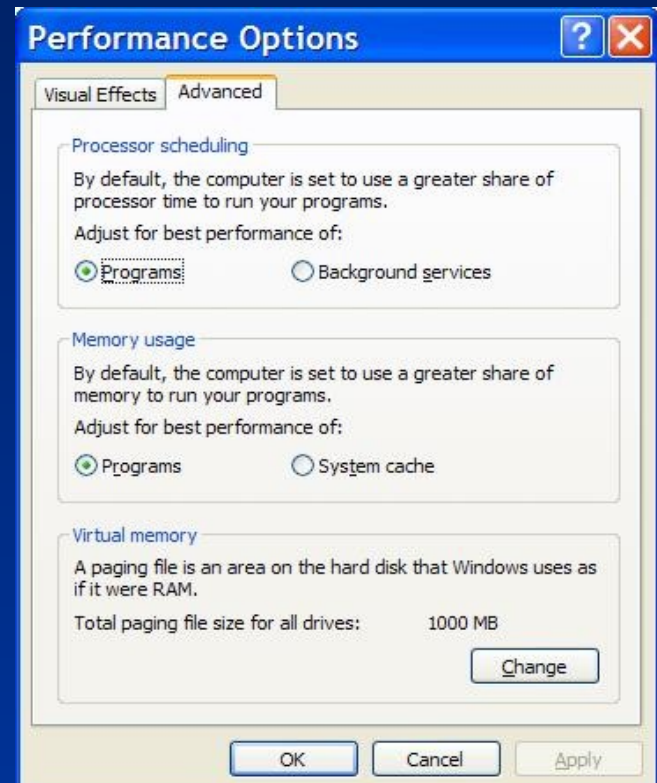
- From Windows 2000, you can choose short or long quantum (e.g. for Terminal Servers)
- NT Server 4.0 was always the same, regardless of slider bar

Windows 2000:



Screen snapshot from:
Control Panel | System | Advanced tab | Performance

Windows XP:



Quantum Control

- Finer grained quantum control can be achieved by modifying the key:
HKLM\System\CurrentControlSet\Control\PriorityControl\Win32PrioritySeparation
 - 6 bit value



- Short vs. Long
 - 0,3 default (short for NT Client versions, long for NT Server versions)
 - 1 long
 - 2 short
- Variable vs. Fixed
 - 0,3 default (yes for NT Client versions, no for NT Server versions)
 - 1 yes
 - 2 no
- Quantum Boost
 - 0 fixed (overrides above setting)
 - 1 double quantum of foreground threads
 - 2,3 triple quantum of foreground threads

Controlling Quantum with Jobs

- If a process is a member of a job, quantum can be adjusted by setting the “Scheduling Class”
 - Only applies if process is higher then Idle priority class
 - Only applies if system running with fixed quanta (the default on NT Servers)
- Values are 0-9
 - 5 is default

Scheduling class	Quantum units
0	6
1	12
2	18
3	24
4	30
5	36
6	42
7	48
8	54
9	60

Priority Boosting

Common boost values
(see NTDDK.H)
1: disk, CD-ROM, parallel,
Video
2: serial, network, named
pipe, mailslot
6: keyboard or mouse
8: sound

- After an I/O: specified by device driver
 - IoCompleteRequest(Irp, PriorityBoost)
- After a **wait on executive event** or **semaphore**
 - Boost value of 1 is used for these objects
 - Server 2003: for critical sections and pushlocks:
 - Waiting thread is boosted to 1 more than setting thread's priority (max boost is to 13)
 - Setting thread loses boost (lock convoy issue)
- After any **wait on a dispatcher object** by a thread in the foreground process:
 - Boost value of 2
 - XP/2003: boost is lost after one full quantum
 - Goal: improve responsiveness of interactive apps
- **GUI threads that wake up** to process windowing input (e.g. windows messages) get a boost of 2
 - This is added to the current, not base priority
 - Goal: improve responsiveness of interactive apps

CPU Starvation Avoidance

- Balance Set Manager system thread looks for “starved” threads
 - This is a thread, running at priority 16
 - Wakes up once per second and examines Ready queues
 - Looks for threads that have been Ready for 300 clock ticks (approximate 4 seconds on a 10ms clock)
 - Attempts to resolve “priority inversions” (example: high priority thread (12 in diagram) waits on something locked by a lower thread (4 in diagram), which can’t run because of a middle priority CPU-bound thread (7 in diagram)), but not deterministically (no priority inheritance)
- Priority is boosted to 15 (14 prior to NT 4 SP3)
 - Quantum is doubled on Windows 2000/XP and set to 4 on 2003
 - At quantum end, returns to previous priority (no gradual decay) and normal quantum
- Scans up to 16 Ready threads per priority level each pass
- Boosts up to 10 Ready threads per pass
- Like all priority boosts, does not apply in the real-time range (priority 16 and above)



Multiprocessor Scheduling

- Threads can run on any CPU, unless specified otherwise
 - Tries to keep threads on same CPU (“soft affinity”)
 - Setting of which CPUs a thread will run on is called “hard affinity”
- Fully distributed (no “master processor”)
 - Any processor can interrupt another processor to schedule a thread
- Scheduling database:
 - Pre-Windows Server 2003: single system-wide list of ready queues
 - Windows Server 2003: per-CPU ready queues

Hard Affinity

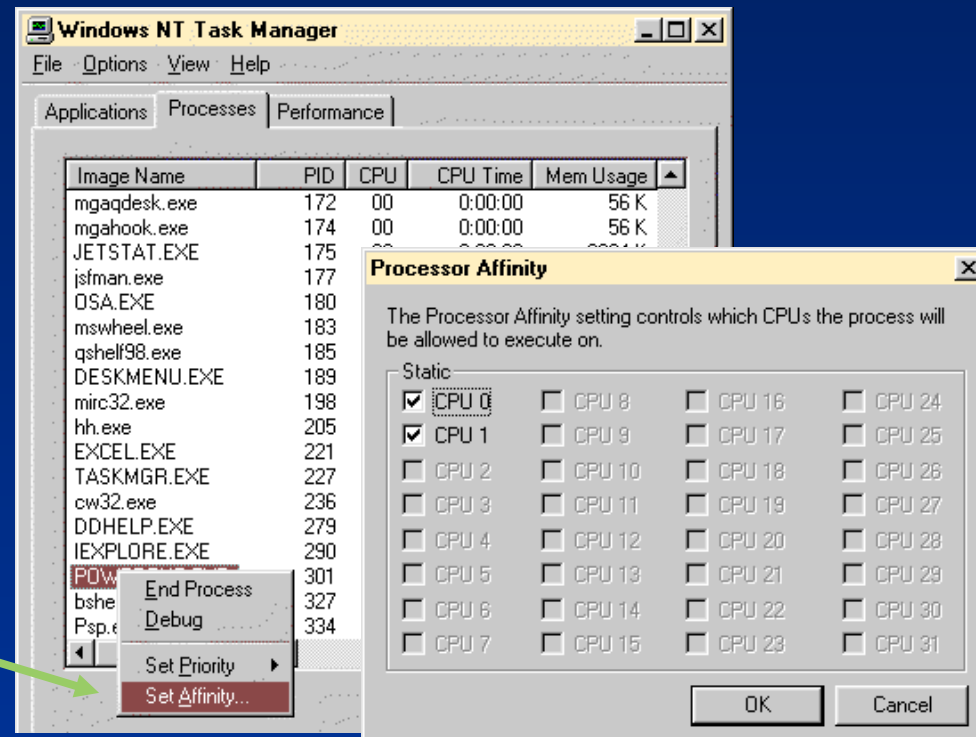
- Affinity is a bit mask where each bit corresponds to a CPU number
 - Hard Affinity specifies where a thread is permitted to run
 - Defaults to all CPUs
 - Thread affinity mask must be subset of process affinity mask, which in turn must be a subset of the active processor mask

- Functions to change:

- SetThreadAffinityMask,*
 - SetProcessAffinityMask,*
 - SetInformationJobObject*

- Tools to change:

- Task Manager or Process Explorer
 - Right click on process and choose "Set Affinity"
 - Psexec -a



Hard Affinity

- Can also set an image affinity mask
 - See “Imagecfg” tool in Windows 2000 Server Resource Kit Supplement 1
 - E.g. `Imagecfg -a 2 xyz.exe` will run xyz on CPU 1
- Can also set “uniprocessor only”: sets affinity mask to one processor
 - `Imagecfg -u xyz.exe`
 - System chooses 1 CPU for the process
 - Rotates round robin at each process creation
 - Useful as temporary workaround for multithreaded synchronization bugs that appear on MP systems
- NOTE: Setting hard affinity can lead to threads’ getting less CPU time than they normally would
 - More applicable to large MP systems running dedicated server apps
 - Also, OS may in some cases run your thread on CPUs other than your hard affinity setting (flushing DPCs, setting system time)
 - Thread “system affinity” vs “user affinity”

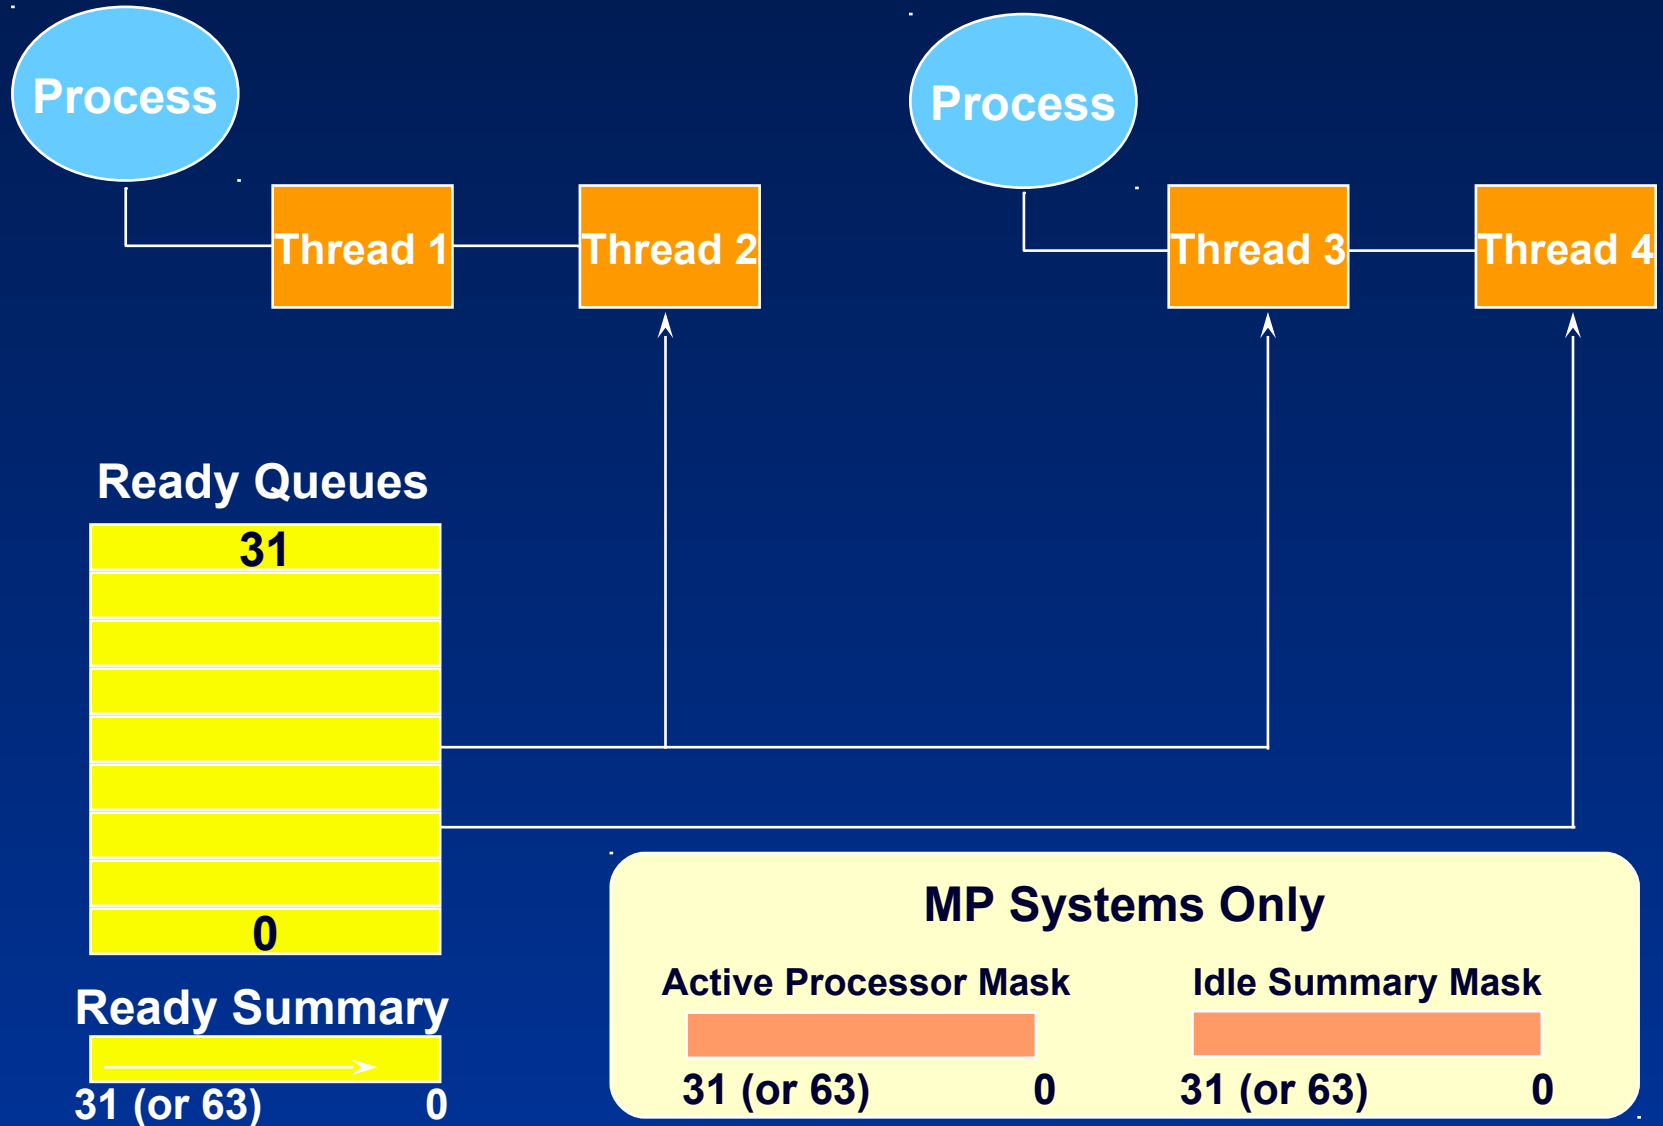
Soft Processor Affinity

- Every thread has an “ideal processor”
 - System selects ideal processor for first thread in process (round robin across CPUs)
 - Next thread gets next CPU relative to the process seed
 - Can override with:

```
SetThreadIdealProcessor (  
    HANDLE hThread,           // handle to the thread to be changed  
    DWORD dwIdealProcessor); // processor number
```

- Hard affinity changes update ideal processor settings
- Used in selecting where a thread runs next (see next slides)
- For Hyperthreaded systems, new Windows API in Server 2003 to allow apps to optimize
 - GetLogicalProcessorInformation
- For NUMA systems, new Windows APIs to allow applications to optimize:
 - Use GetProcessAffinityMask to get list of processors
 - Then GetNumaProcessorNode to get node # for each CPU
 - Or call GetNumaHighestNodeNumber and then GetNumaNodeProcessorMask to get the processor #s for each node

Windows 2000/XP Dispatcher Database



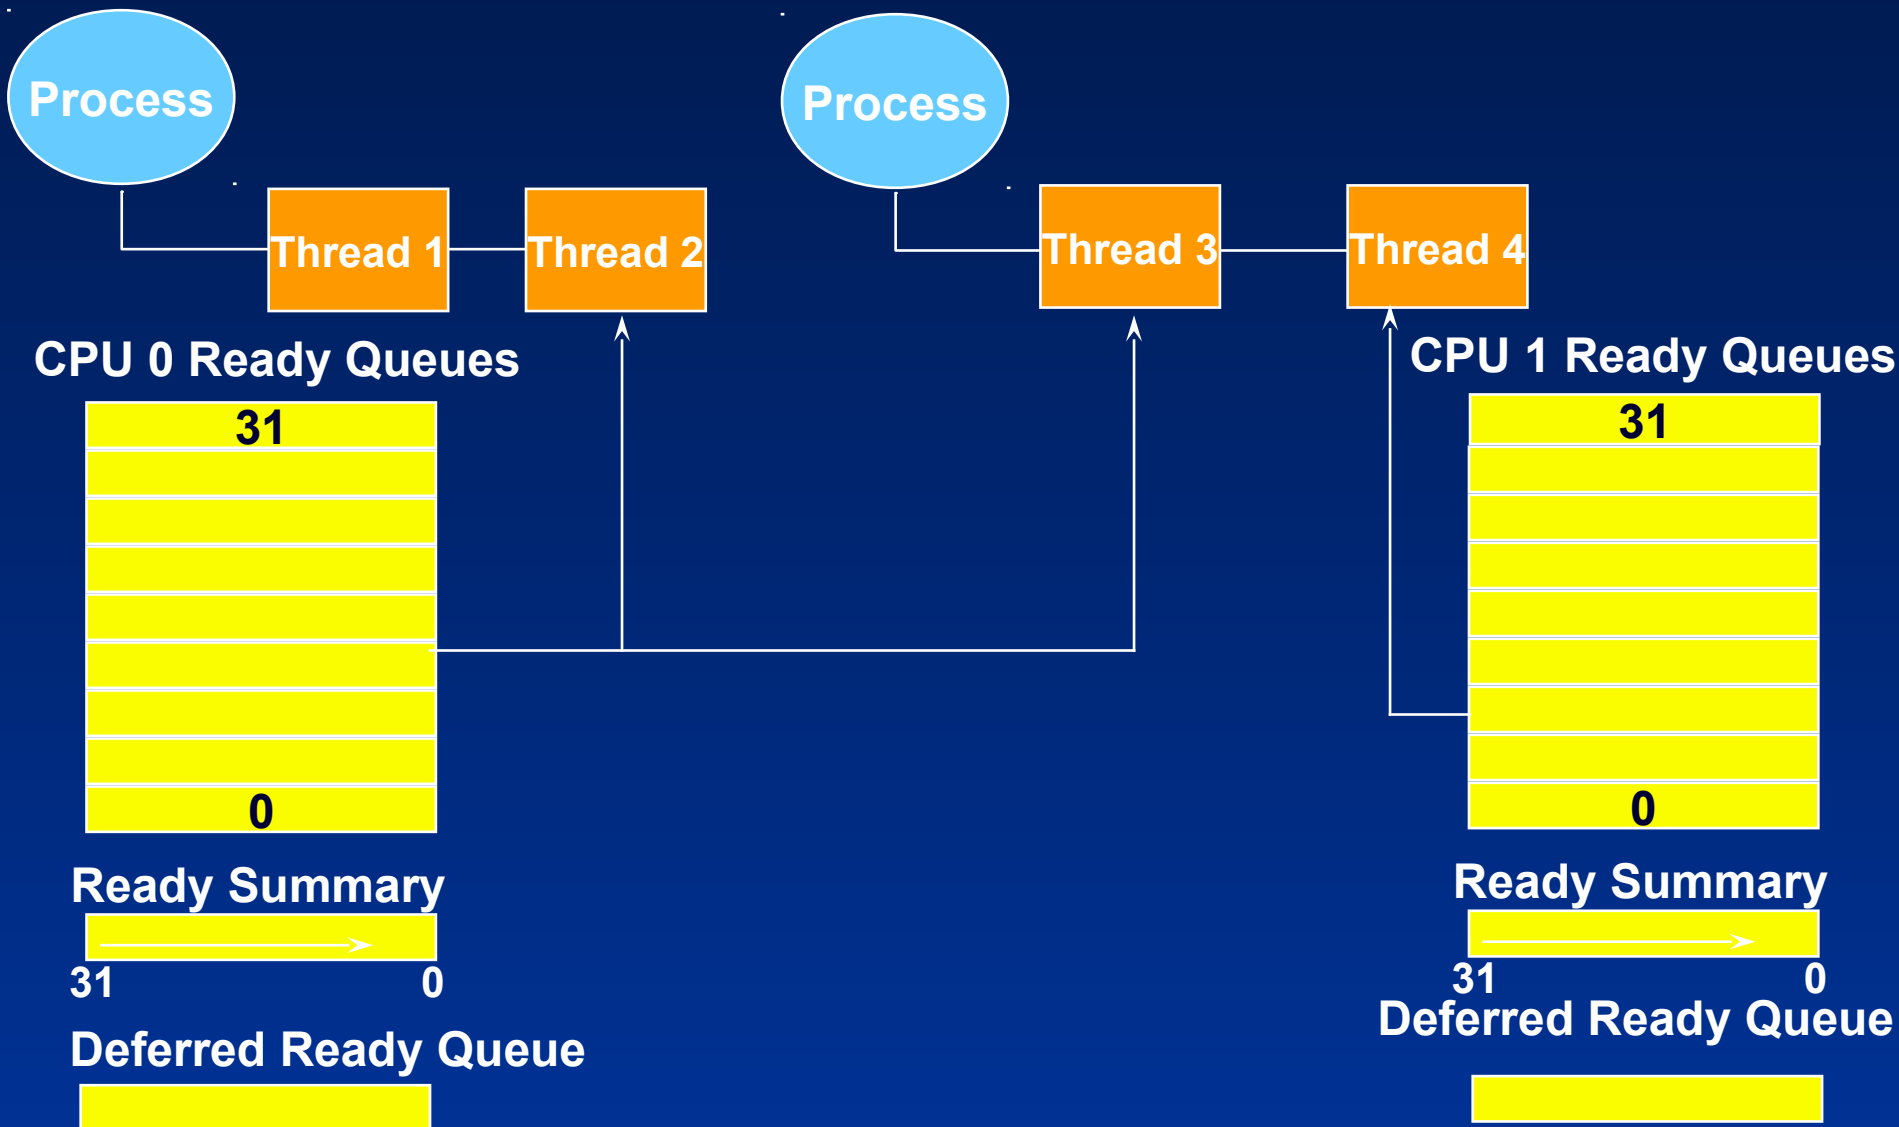
Choosing a CPU for a Ready Thread (Windows 2000 & XP)

- When a thread becomes ready to run (e.g. its wait completes, or it is just beginning execution), need to choose a processor for it to run on
- First, it sees if any processors are idle that are in the thread's hard affinity mask:
 - If its "ideal processor" is idle, it runs there
 - Else, if the previous processor it ran on is idle, it runs there
 - Else if the current processor is idle, it runs there
 - Else it picks the highest numbered idle processor in the thread's affinity mask
- If no processors are idle:
 - If the ideal processor is in the thread's affinity mask, it selects that
 - Else if the the last processor it ran on is in the thread's affinity mask, it selects that
 - Else it picks the highest numbered processor in the thread's affinity mask
- Finally, it compares the priority of the new thread with the priority of the thread running on the processor it selected (if any) to determine whether or not to perform a preemption

Selecting a Thread to Run on a CPU (Windows 2000 & XP)

- System needs to choose a thread to run on a specific CPU in the following cases:
 - At quantum end
 - When a thread enters a wait state
 - When a thread removes its current processor from its hard affinity mask
 - When a thread exits
- Starting with the first thread in the highest priority non-empty ready queue, it scans the queue for the first thread that has the current processor in its hard affinity mask and:
 - Ran last on the current processor, or
 - Has its ideal processor equal to the current processor, or
 - Has been in its Ready queue for 3 or more clock ticks, or
 - Has a priority ≥ 24
- If it cannot find such a candidate, it selects the highest priority thread that can run on the current CPU (whose hard affinity includes the current CPU)
 - Note: this may mean going to a lower priority ready queue to find a candidate

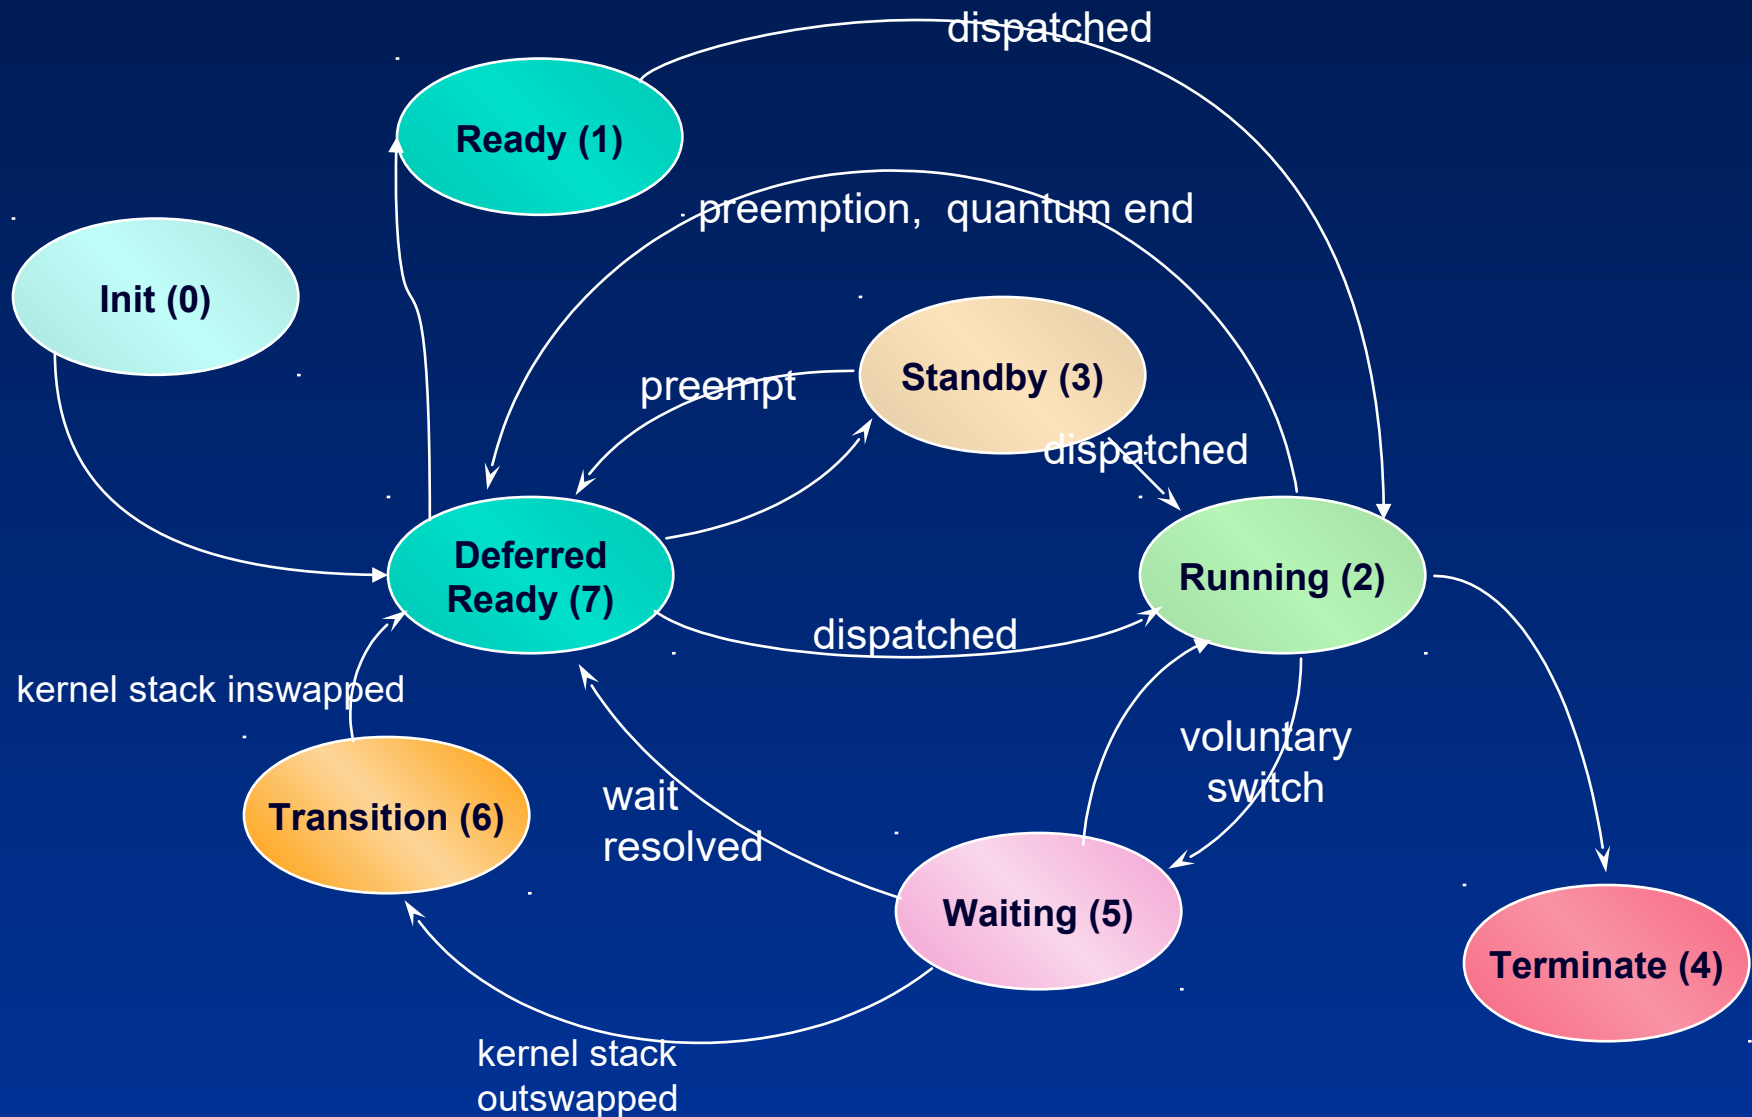
Windows Server2003 Dispatcher Database



Server 2003 Enhancements

- Threads always go into the ready queue of their ideal processor
- Instead of locking the dispatcher database to look for a candidate to run, per-CPU ready queue is checked first (first grabs PRCB spinlock)
 - If a thread has been selected to run on the CPU, does the context switch
 - Else begins scan of other CPU's ready queues looking for a thread to run
 - This scan is done OUTSIDE the dispatcher lock
 - Just acquires CPU PRCB lock
- Dispatcher lock still acquired to wait or unwait a thread and/or change state of a dispatcher object
- Bottom line: dispatcher lock is now held for a MUCH shorter time

Thread Scheduling States (Server 2003)

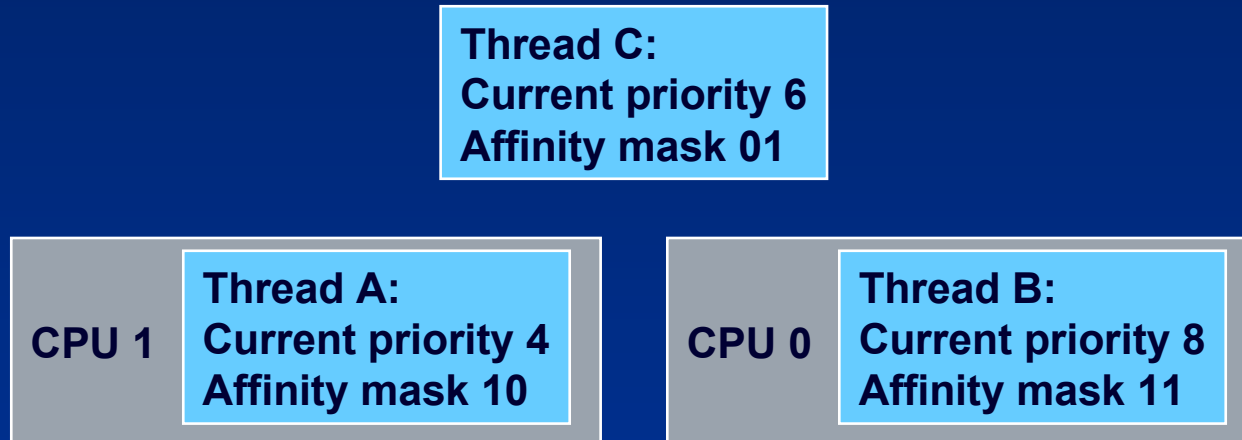


Server 2003 Enhancements

- Idle processor selection further refined to:
 - If a NUMA system: if there are idle CPUs in the node containing the thread's ideal processor, reduce to that set
 - If a hyperthreaded system: if one of the idle processors is a physical processor with all logical processors idle, reduce to that set
 - Then try to eliminate idle CPUs that are sleeping
 - If thread ran last on a member of the set, pick that CPU
 - Else pick lowest numbered CPU in remaining set

“Affinity Collisions”

- Highest-priority n threads may not be Running if thread affinity interferes
- NT guarantees the highest-priority thread will be Running
 - But lower-priority $n-1$ Ready threads may not be...
 - ...because scheduler will not “move” running threads among CPUs
- Example: Threads became Ready in order A, B, C



Further Reading

- Pavel Yosifovich, Alex Ionescu, et al., “Windows Internals”, 7th Edition, Microsoft Press, 2017.
 - Chapter 4 – Threads (from pp. 275)
 - Thread scheduling (from pp. 303)
 - Scheduling scenarios (from pp. 361)
 - Multiprocessor systems (from pp. 376)