

Unit 5: Memory Management

5.2. Windows Memory Management Fundamentals

Roadmap for Section 5.2.

- Memory Manager Features and Components
- Virtual Address Space Allocation
- Shared Memory and Memory-Mapped Files
- Physical Memory Limits
- Memory management APIs

Windows Memory Management Fundamentals

- Classical virtual memory management
 - Flat virtual address space per process
 - Private process address space
 - Global system address space
 - Per session address space
- Object based
 - Section object and object-based security (ACLs...)
- Demand paged virtual memory
 - Pages are read in on demand & written out when necessary (to make room for other memory needs)
- Provides flat virtual address space
 - 32-bit: 4 GB, 64-bit: 16 Exabytes (theoretical)

Windows Memory Management Fundamentals

- Lazy evaluation
 - Sharing – usage of prototype PTEs (page table entries)
 - Extensive usage of `copy_on_write`
 - ...whenever possible
- Shared memory with copy on write
- Mapped files (fundamental primitive)
 - Provides basic support for file system cache manager

Memory Manager Components

- System services for allocating, deallocating, and managing virtual memory
- A access fault trap handler for resolving hardware-detected memory management exceptions and making virtual pages resident on behalf of a process
- Six system threads
 - *Working set manager* (priority 16) – drives overall memory management policies, such as working set trimming, aging, and modified page writing
 - *Process/stack swapper* (priority 23) -- performs both process and kernel thread stack inswapping and outswapping
 - *Modified page writer* (priority 17) – writes dirty pages on the modified list back to the appropriate paging files
 - *Mapped page writer* (priority 17) – writes dirty pages from mapped files to disk
 - *Dereference segment thread* (priority 18) is responsible for cache and page file growth and shrinkage
 - *Zero page thread* (priority 0) – zeros out pages on the free list

MM: Process Support

- MmCreateProcessAddressSpace – 3 pages
 - The page directory
 - Points to itself
 - Map the page table of the hyperspace
 - Map system paged and nonpaged areas
 - Map system cache page table pages
 - The page table page for working set
 - The page for the working set list
- MmInitializeProcessAddressSpace
 - Initialize PFN for PD and hyperspace PDEs
 - MiInitializeWorkingSetList
 - Optional: MmMapViewOfSection for image file
- MmCleanProcessAddressSpace
- MmDeleteProcess AddressSpace

MM: Process Swap Support

- MmOutSwapProcess / MmInSwapProcess
- MmCreateKernelStack
 - MiReserveSystemPtes for stack and no-access page
- MmDeleteKernelStack
 - MiReleaseSystemPtes
- MmGrowKernelStack
- MmOutPageKernelStack
 - Signature (thread_id) written on top of stack before write
 - The page goes to transition list
- MmInPageKernelStack
 - Check signature after stack page is read / bugcheck

MM: Working Sets

Working Set:

- The set of pages in memory at any time for a given process, or
- All the pages the process can reference without incurring a page fault
- Per process, private address space
- WS limit: maximum amount of pages a process can own
- Implemented as array of working set list entries (WSLE)

Soft vs. Hard Page Faults:

- Soft page faults resolved from memory (standby/modified page lists)
- Hard page faults require disk access

Working Set Dynamics:

- Page replacement when WS limit is reached
- NT 4.0: page replacement based on modified FIFO
- Windows 2000: Least Recently Used algorithm (uniprocessor systems)

MM: Working Set Management

- Modified Page Writer thread
 - Created at system initialization
 - Writing modified pages to backing file
 - Optimization: min. I/Os, contiguous pages on disk
 - Generally MPW is invoked before trimming
- Balance Set Manager thread
 - Created at system initialization
 - Wakes up every second
 - Executes MmWorkingSetManager
 - Trimming process WS when required: from current down to minimal WS for processes with lowest page fault rate
 - Aware of the system cache working set
 - Process can be out-swapped if all threads have pageable kernel stack

MM: I/O Support

- I/O Support operations:
 - Locking/Unlocking pages in memory
 - Mapping/Unmapping Locked Pages into current address space
 - Mapping/Unmapping I/O space
 - Get physical address of a locked page
 - Probe page for access
- Memory Descriptor List
 - Starting VAD
 - Size in Bytes
 - Array of elements to be filled with physical page numbers
- Physically contiguous vs. Virtually contiguous

MM: Cache Support

- System wide cache memory
 - Region of system paged area reserved at initialization time
 - Initial default: 512 MB (min. 64MB if /3GB, max. 960 MB)
 - Managed as system wide working set
 - A valid cache page is valid in all address spaces
 - Lock the page in the cache to prevent WS removal
 - WS Manager trimming thread is aware of this special WS
 - Not accessible from user mode
 - Only views of mapped files may reside in the cache
- File Systems and Server interaction support
 - Map/Unmap view of section in system cache
 - Lock/Unlock pages in system cache
 - Read section file in system cache
 - Purge section

Memory Manager: Services

- Caller can manipulate own/remote memory
 - Parent process can allocate/deallocate, read/write memory of child process
 - Subsystems manage memory of their client processes this way
- Most services are exposed through Windows API
 - Page granularity virtual memory functions (Virtualxxx...)
 - Memory-mapped file functions (CreateFileMapping, MapViewOfFile)
 - Heap functions (Heapxxx...)
- Services for device drivers/kernel code (Mm...)

Protecting Memory

Attribute	Description
PAGE_NOACCESS	Read/write/execute causes access violation
PAGE_READONLY	Write/execute causes access violation; read permitted
PAGE_READWRITE	Read/write accesses permitted
PAGE_EXECUTE	Any read/write causes access violation; execution of code is permitted (relies on special processor support)
PAGE_EXECUTE_READ	Read/execute access permitted (relies on special processor support)
PAGE_EXECUTE_READWRITE	All accesses permitted (relies on special processor support)
PAGE_WRITECOPY	Write access causes the system to give process a private copy of this page; attempts to execute code cause access violation
PAGE_EXECUTE_WRITECOPY	Write access causes creation of private copy of page
PAGE_GUARD	Any read/write attempt raises EXCEPTION_GUARD_PAGE and turns off guard page status

Reserving & Committing Memory

- Optional 2-phase approach to memory allocation:
 1. Reserve address space (in multiples of page size)
 2. Commit storage in that address space
 - Can be combined in one call (`VirtualAlloc`, `VirtualAllocEx`)
- Reserved memory:
 - Range of virtual addresses reserved for future use (contiguous buffer)
 - Accessing reserved memory results in access violation
 - Fast, inexpensive
- Committed memory:
 - Has backing store (pagefile.sys, memory-mapped file)
 - Either private or mapped into a view of a section
 - Decommit via `VirtualFree`, `VirtualFreeEx`

A thread's user-mode stack is constructed using this 2-phase approach: initial reserved size is 1MB, only 2 pages are committed: stack & guard page

Features new to Windows 2000

Memory Management

- Support of 64 GB physical memory on Intel platform
 - PAE – Physical Address Extension (36 bit, changes PDE/PTE structs)
 - New version of kernel (ntkrnlpa.exe, ntkrpamp.exe)
 - /PAE switch in boot.ini
- Integrated support for Terminal Server
 - HydraSpace : per session
 - In NT 4.0 Terminal Server had a specific kernel
- Driver Verifier: verifier.exe
 - Pool checking, IRQL checking
 - Low resources simulation, pool tracking, I/O verification

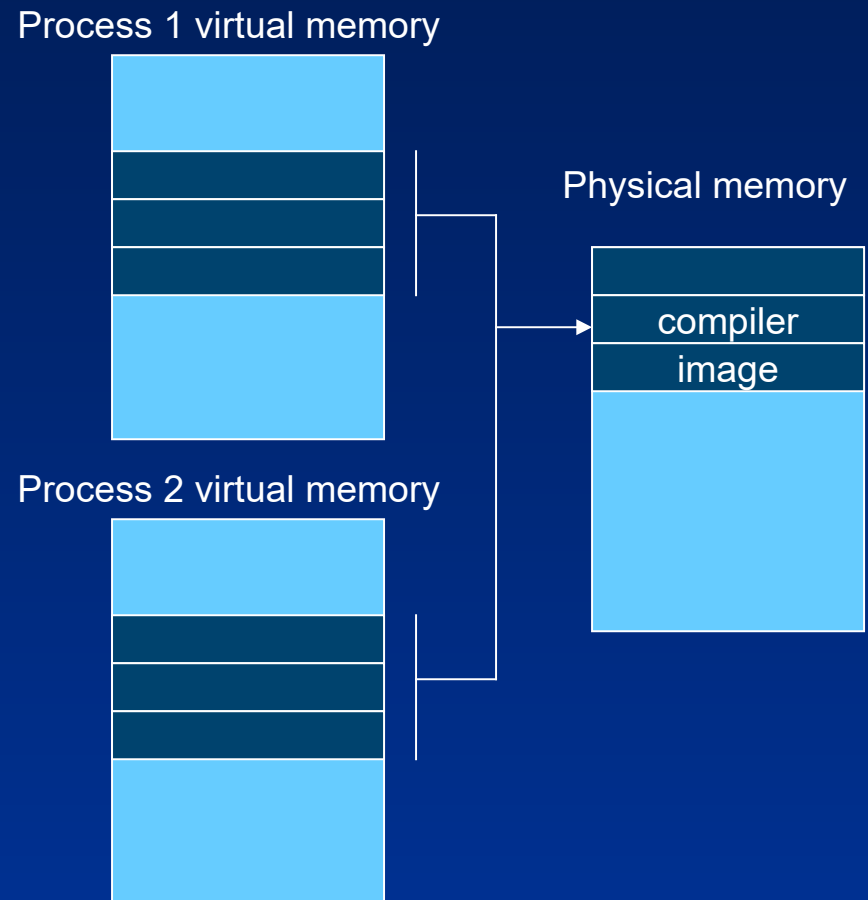
Features new to Windows XP/2003

Memory Management

- 64-bit support
- Up to 1024 Gbytes physical memory supported
- Support for Data Execution Prevention (DEP)
 - Memory manager supports HW no-execute protection
- Performance & Scalability enhancements

Shared Memory & Mapped Files

- Shared memory + copy-on-write per default
- Executables are mapped as read-only
- Memory manager uses section objects to implement shared memory (file mapping objects in Windows API)



Virtual Address Space Allocation

- Virtual address space is sparse
 - Address spaces contain reserved, committed, and unused regions
- Unit of protection and usage is one page
 - On x86, default page size is 4 KB (x86 supports 4KB and 4MB)
 - In PAE mode, large pages are 2 MB
 - On x64, default page size is 4 KB; large pages are 4 MB
 - On Itanium, default page size is 8 KB; large pages are 16 MB (Itanium supports 4KB, 8KB, 16KB, 64KB, 256KB, 1MB, 4MB, 16MB, 64MB and 256MB)

Large Pages

- Large pages allow a single page directory entry to map a larger region
 - x86, x64: 4 MB, IA64: 16 MB
 - Advantage: improves performance
 - Single TLB entry used to map larger area
- Large pages are used to map NTOSKRNL, HAL, nonpaged pool, and the PFN database if a “large memory system”
 - Windows 2000: more than 127 MB
 - Windows XP/2003: more than 255 MB
 - In other words, most systems...
- Disadvantage: disables kernel write protection
 - With small pages, OS/driver code pages are mapped as read only; with large pages, entire area must be mapped read/write
 - Drivers can then modify/corrupt system & driver code without immediately crashing system
 - Driver Verifier turns large pages off
 - Can also override by changing
HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Control\Session Manager\Memory Management\LargePageMinimum to FFFFFFFF

Large Pages: Server 2003 Enhancements

- Can specify other drivers to map with large pages:
 - HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Control\Session Manager\Memory Management\LargePageDrivers (multi-string)
- Applications can use large pages for process memory
 - VirtualAlloc with MEM_LARGE_PAGE flag
 - Can query if system supports large pages with GetLargePageMinimum

Data Execution Prevention

- Windows XP SP2 and Windows Server 2003 SP1 support Data Execution Prevention (DEP)
 - Prevents code from executing in a memory page not specifically marked as executable
 - Stops exploits that rely on getting code executed in data areas
- Relies on hardware ability to mark pages as non executable
 - AMD calls it NX (“No Execute”)
 - Intel calls it XD (“Execute Disable”)
- Processor support:
 - Intel Itanium had this in 2001, but Windows didn’t support it until XP SP2 / Server 2003 SP1
 - AMD64 was the next to support it
 - Then, AMD added Sempron (32-bit processor with NX support)
 - Intel added it first with their 64-bit extension chips (Xeon/Pentium4 with EM64T)
 - Then, Intel added it to their 32-bit processor line (anything ending in “J”)

Data Execution Prevention

- Attempts to execute code in a page marked no execute result in:
 - User mode: access violation exception
 - Kernel mode: `ATTEMPTED_EXECUTE_OF_NOEXECUTE_MEMORY` bugcheck (blue screen)
- Memory that needs to be executable must be marked as such using page protection bits on `VirtualAlloc` and `VirtualProtect` APIs:
 - `PAGE_EXECUTE`, `PAGE_EXECUTE_READ`,
`PAGE_EXECUTE_READWRITE`, `PAGE_EXECUTE_WRITECOPY`

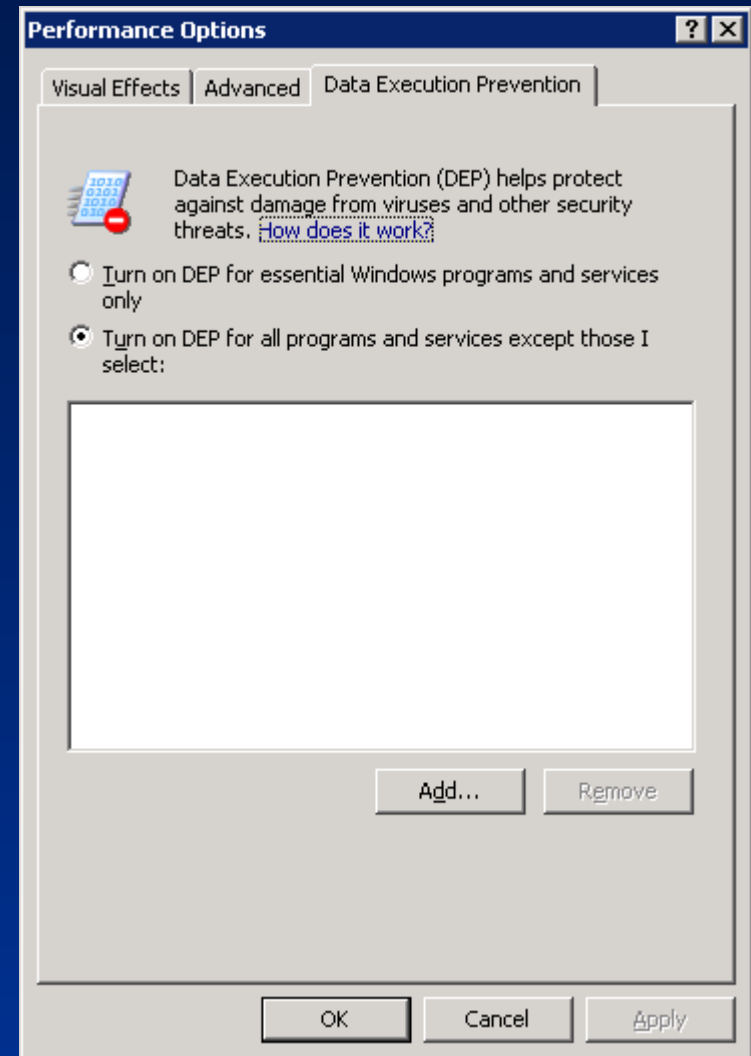
Controlling DEP

- New Boot.ini switch /NOEXECUTE
 - /NOEXECUTE=ALWAYSON – enables DEP for all applications
 - /NOEXECUTE=ALWAYSOFF – disables DEP
- Two qualifiers apply only to 32-bit applications:
 - /NOEXECUTE=OPTIN – enables DEP for core Windows programs
 - Default for Windows XP (32-bit and 64-bit editions)
 - /NOEXECUTE=OPTOUT – enables DEP for all applications except those excluded
 - Default for Windows Server 2003 (32-bit and 64-bit editions)
- Control Panel setting (in all Windows versions, incl. Windows 10):

Control Panel → System → 'Advanced system settings' → 'Performance' settings → DEP settings

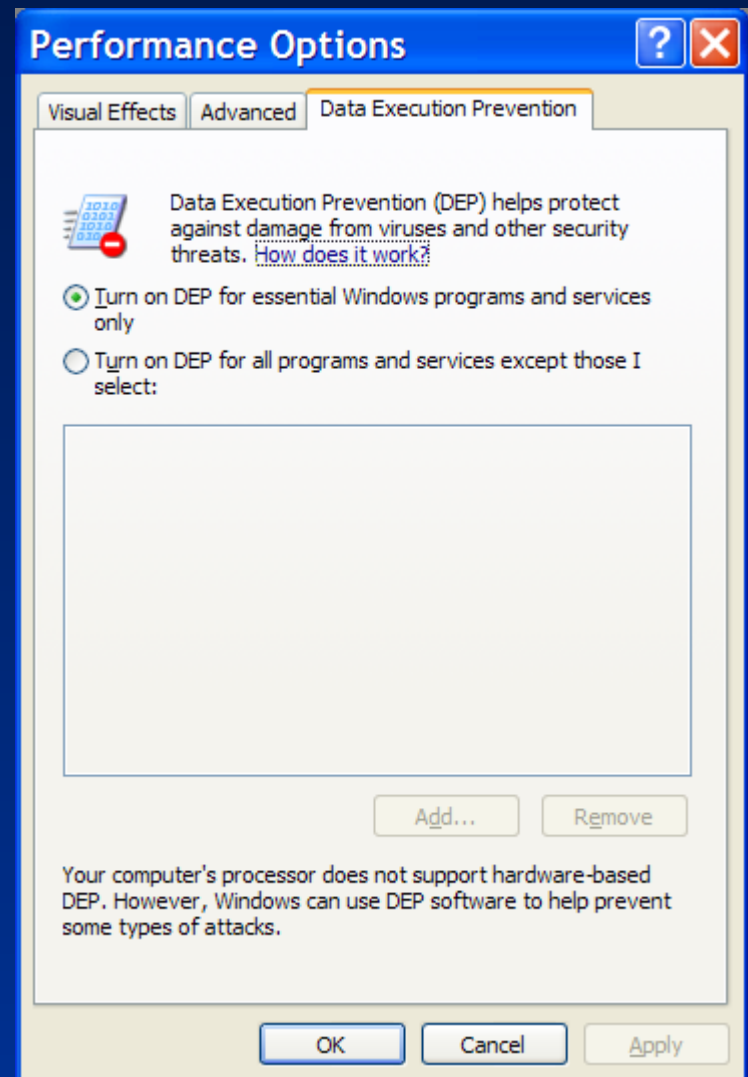
DEP on 64-bit Windows

- Always applied to all 64-bit processes and device drivers
 - Protects user and kernel stacks, paged pool, session pool
- 32-bit processes depend on configuration settings



DEP on 32-bit Windows

- Hardware DEP used when running 32-bit Windows on systems that support it
- When enabled, system boots PAE kernel (Ntkrnlpa.exe)
- Kernel mode: applied to kernel stacks, but not paged/session pool
- User mode: depends on system configuration
- Even on processors without hardware DEP, some limited protection implemented for exception handlers



Mapped Files

- A way to take part of a file and map it to a range of virtual addresses (address space is 2 GB, but files can be much larger)
- Called “file mapping objects” in Windows API
- Bytes in the file then correspond one-for-one with bytes in the region of virtual address space
 - Read from the “memory” fetches data from the file
 - Pages are kept in physical memory as needed
 - Changes to the memory are eventually written back to the file (can request explicit flush)
- Initial mapped files in a process include:
 - The executable image (EXE)
 - One or more Dynamically Linked Libraries (DLLs)
- Processes can map additional files as desired (data files or additional DLLs)



Section Objects (mapped files)

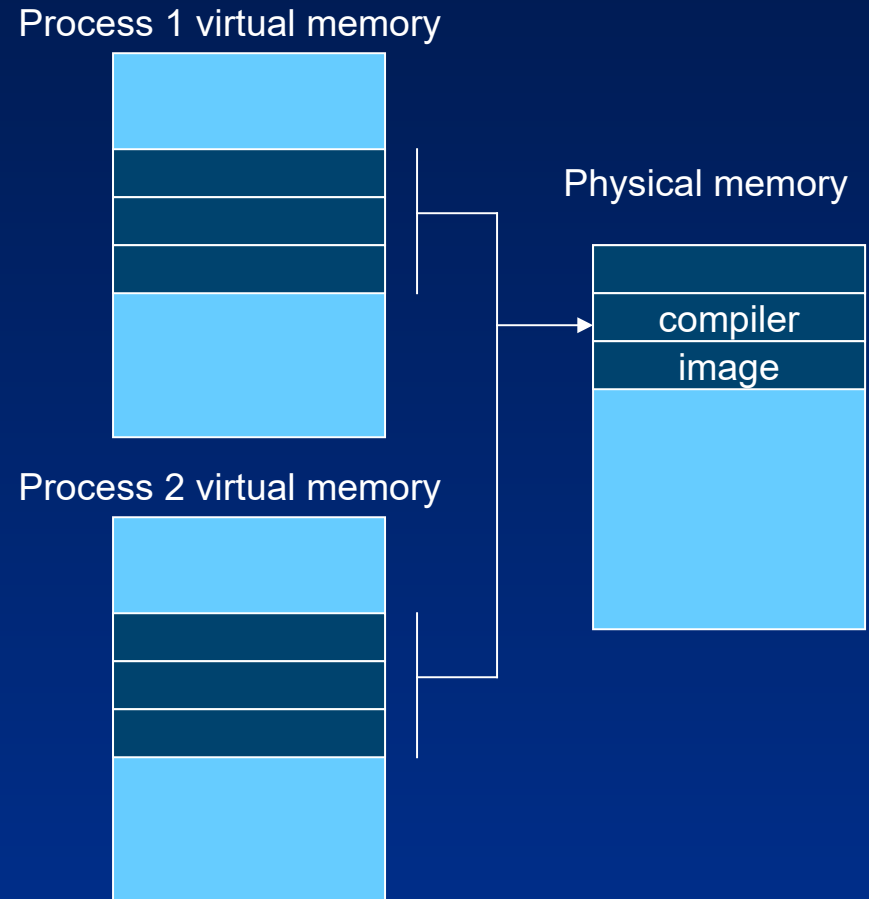
- Called “file mapping objects” in Windows API
- Files may be mapped into virtual address space

```
// first, do EITHER ...  
hMapObj = CreateFileMapping (hFile, security, protection, sizeHigh, sizeLow, mapname);  
// ... OR ...  
hMapObj = OpenFileMapping (accessMode, inheritflag, mapname);  
  
// ... then, pass the resulting handle to a mapping object (section) to ...  
lpvoid = MapViewOfFile (hMapObj, accessMode, offsetHigh, offsetLow, cbMap);
```

- Bytes in the file then correspond one-for-one with bytes in the region of virtual address space
 - Read from the “memory” fetches data from the file
 - Changes to the memory are written back to the file
 - Pages are kept in physical memory as needed
 - If desired, can map to only a part of the file at a time

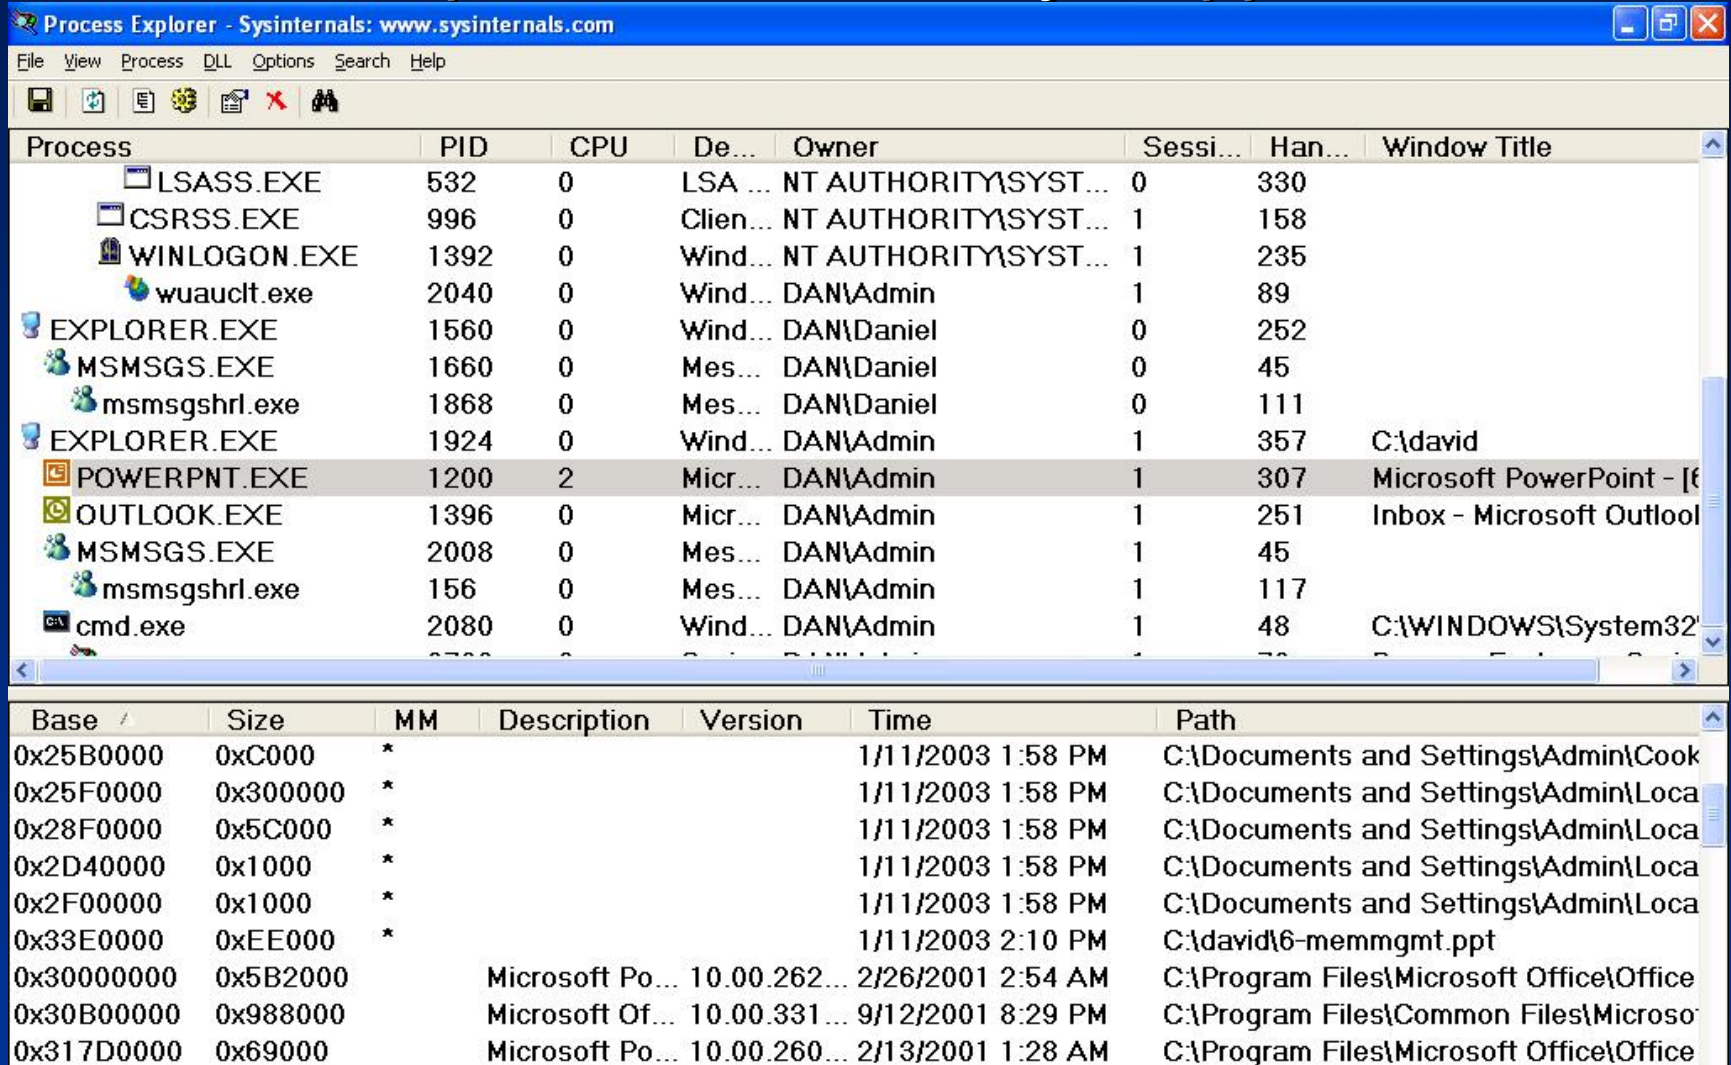
Shared Memory

- Like most modern OS's, Windows provides a way for processes to share memory
 - High speed IPC (used by LPC, which is used by RPC)
 - Threads share address space, but applications may be divided into multiple processes for stability reasons
- It does this automatically for shareable pages
 - E.g. code pages in an EXE or DLL
- Processes can also create shared memory sections
 - Called page file backed file mapping objects
 - Full Windows security



Viewing DLLs & Memory Mapped Files

- Process Explorer lists memory mapped files



The screenshot shows the Process Explorer application window. The top pane displays a list of running processes. The bottom pane shows the memory mapped files for the selected process, 'POWERPNT.EXE'.

Process	PID	CPU	De...	Owner	Sessi...	Han...	Window Title
LSASS.EXE	532	0	LSA ...	NT AUTHORITY\SYST...	0	330	
CSRSS.EXE	996	0	Clie...	NT AUTHORITY\SYST...	1	158	
WINLOGON.EXE	1392	0	Wind...	NT AUTHORITY\SYST...	1	235	
wuauclt.exe	2040	0	Wind...	DAN\Admin	1	89	
EXPLORER.EXE	1560	0	Wind...	DAN\Daniel	0	252	
MSMSG.S.EXE	1660	0	Mes...	DAN\Daniel	0	45	
msmsgshrl.exe	1868	0	Mes...	DAN\Daniel	0	111	
EXPLORER.EXE	1924	0	Wind...	DAN\Admin	1	357	C:\david
POWERPNT.EXE	1200	2	Micr...	DAN\Admin	1	307	Microsoft PowerPoint - [6...
OUTLOOK.EXE	1396	0	Micr...	DAN\Admin	1	251	Inbox - Microsoft Outlool
MSMSG.S.EXE	2008	0	Mes...	DAN\Admin	1	45	
msmsgshrl.exe	156	0	Mes...	DAN\Admin	1	117	
cmd.exe	2080	0	Wind...	DAN\Admin	1	48	C:\WINDOWS\System32\...

Base	Size	MM	Description	Version	Time	Path
0x25B0000	0xC000	*			1/11/2003 1:58 PM	C:\Documents and Settings\Admin\Cook
0x25F0000	0x300000	*			1/11/2003 1:58 PM	C:\Documents and Settings\Admin\Loca
0x28F0000	0x5C000	*			1/11/2003 1:58 PM	C:\Documents and Settings\Admin\Loca
0x2D40000	0x1000	*			1/11/2003 1:58 PM	C:\Documents and Settings\Admin\Loca
0x2F00000	0x1000	*			1/11/2003 1:58 PM	C:\Documents and Settings\Admin\Loca
0x33E0000	0xEE000	*			1/11/2003 2:10 PM	C:\david\6-memmgmt.ppt
0x30000000	0x5B2000		Microsoft Po...	10.00.262...	2/26/2001 2:54 AM	C:\Program Files\Microsoft Office\Office
0x30B00000	0x988000		Microsoft Of...	10.00.331...	9/12/2001 8:29 PM	C:\Program Files\Common Files\Microso
0x317D0000	0x69000		Microsoft Po...	10.00.260...	2/13/2001 1:28 AM	C:\Program Files\Microsoft Office\Office

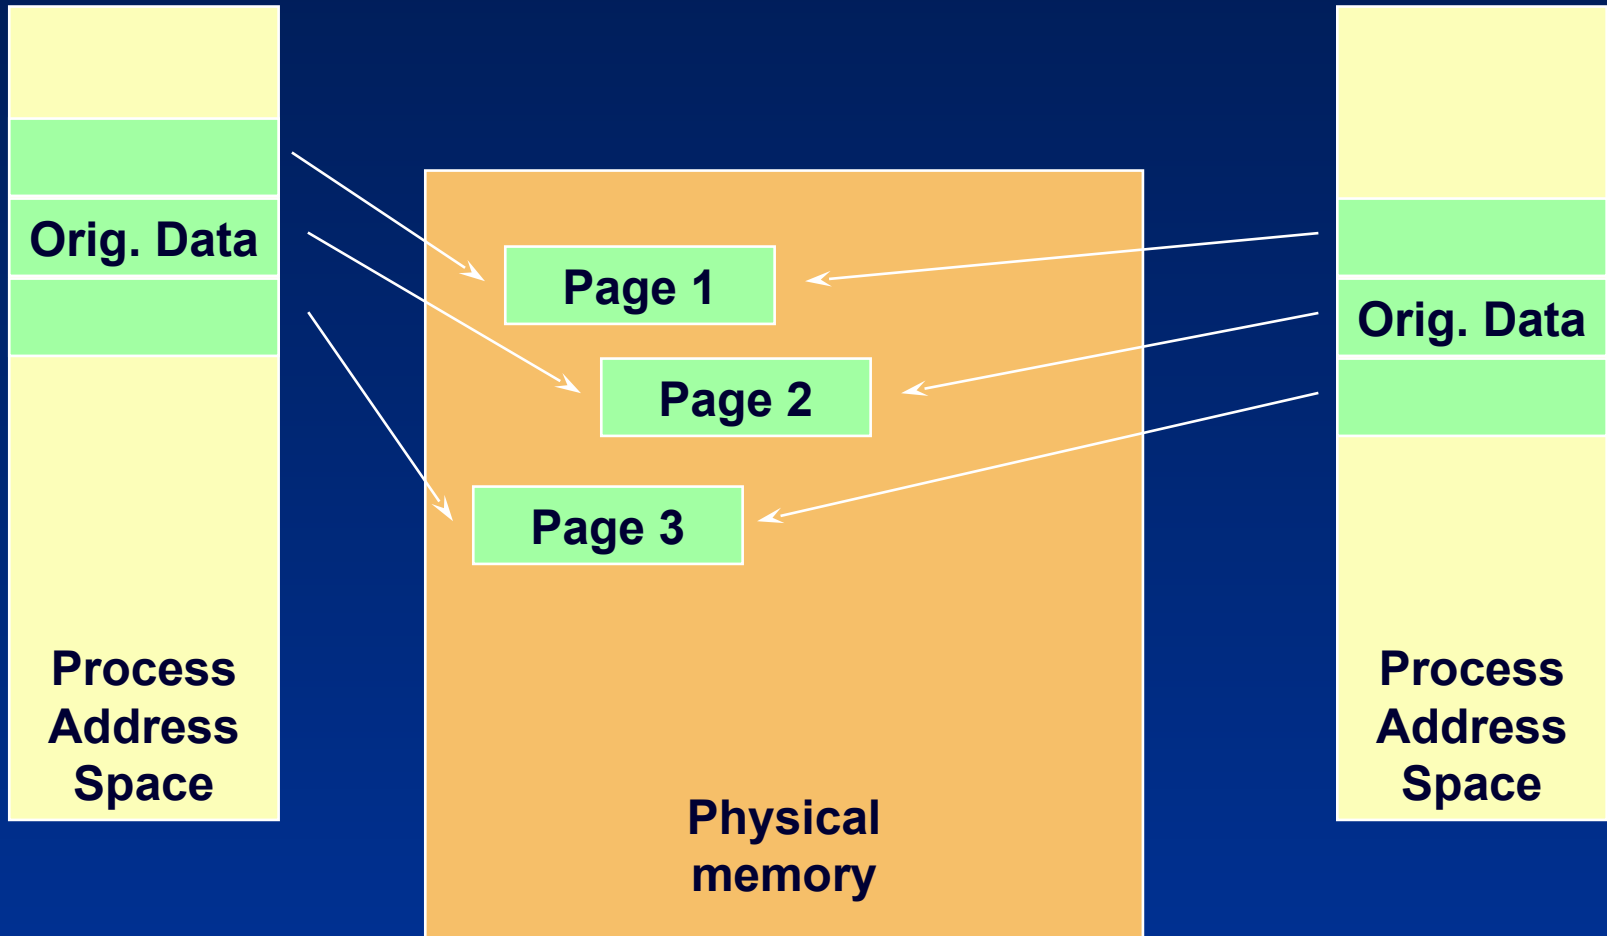


Copy-On-Write Pages

- Used for sharing between process address spaces
- Pages are originally set up as shared, read-only, faulted from the common file
 - Access violation on write attempt alerts pager
 - pager makes a copy of the page and allocates it privately to the process doing the write, backed to the paging file
 - So, only need unique copies for the pages in the shared region that are actually written (example of “lazy evaluation”)
 - Original values of data are still shared
 - e.g. writeable data initialized with C initializers



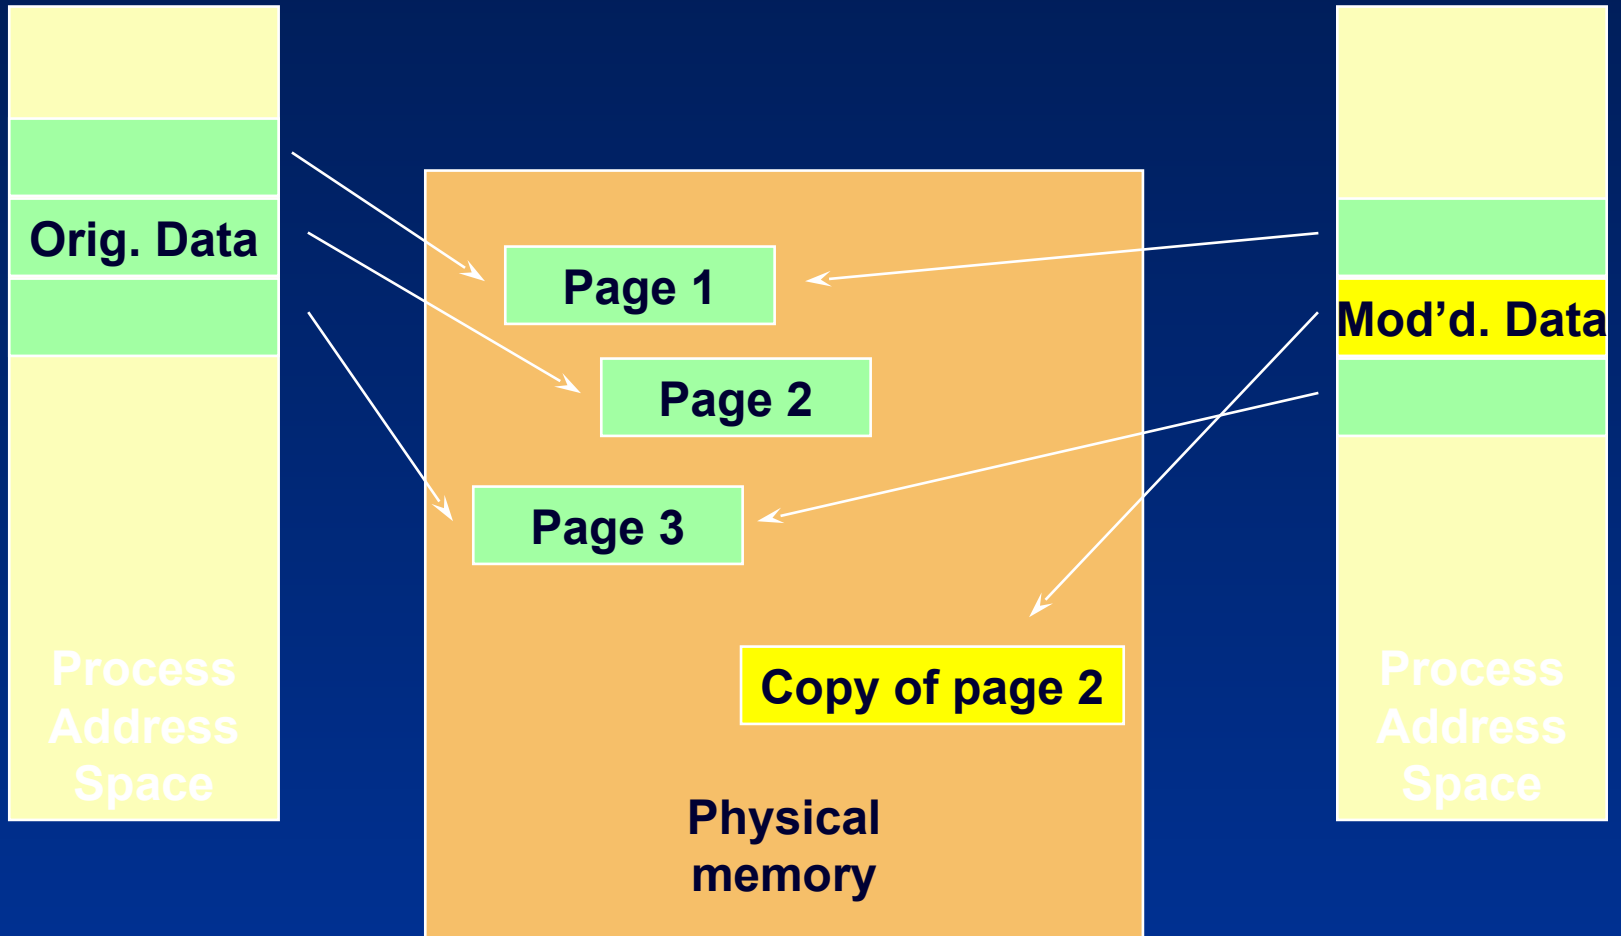
How Copy-On-Write Works Before



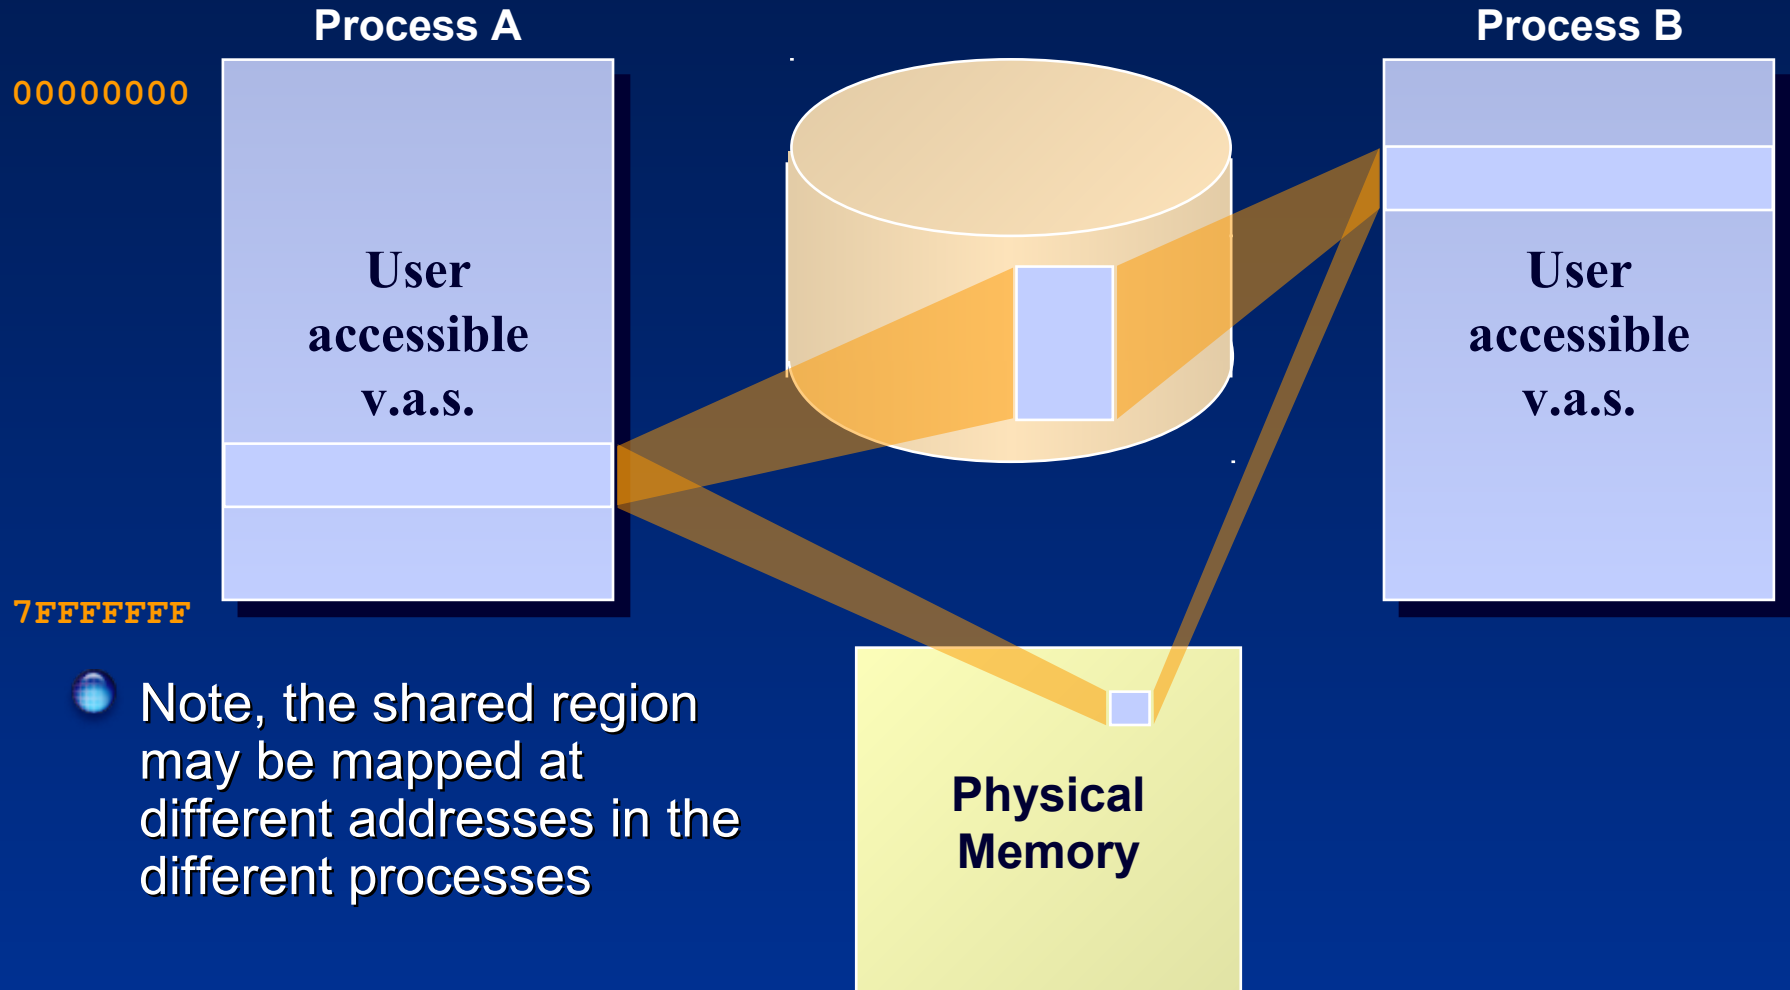


How Copy-On-Write Works

After



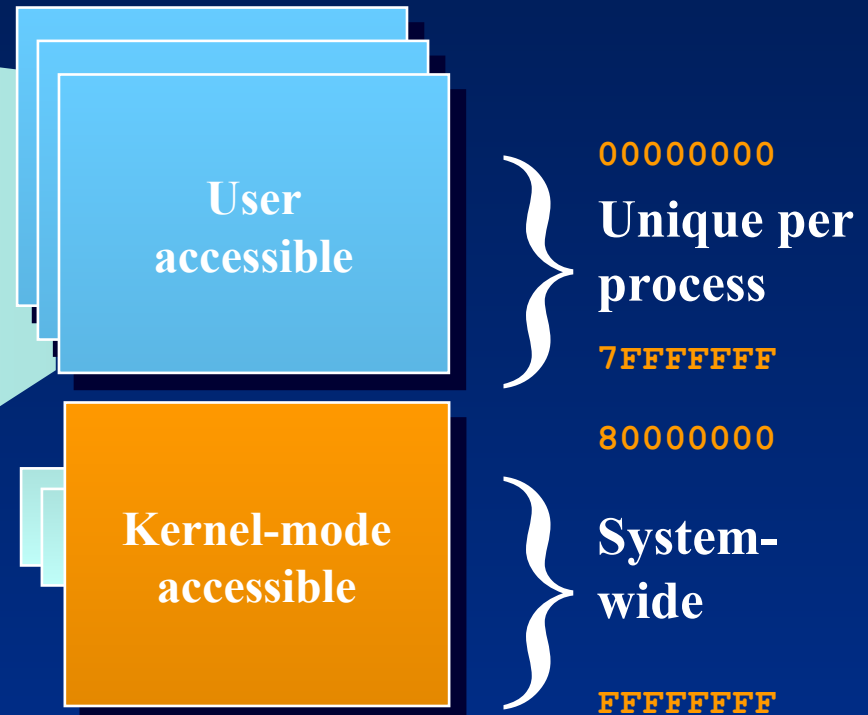
Shared Memory = File Mapped by Multiple Processes



Virtual Address Space (V.A.S.)

Process space contains:

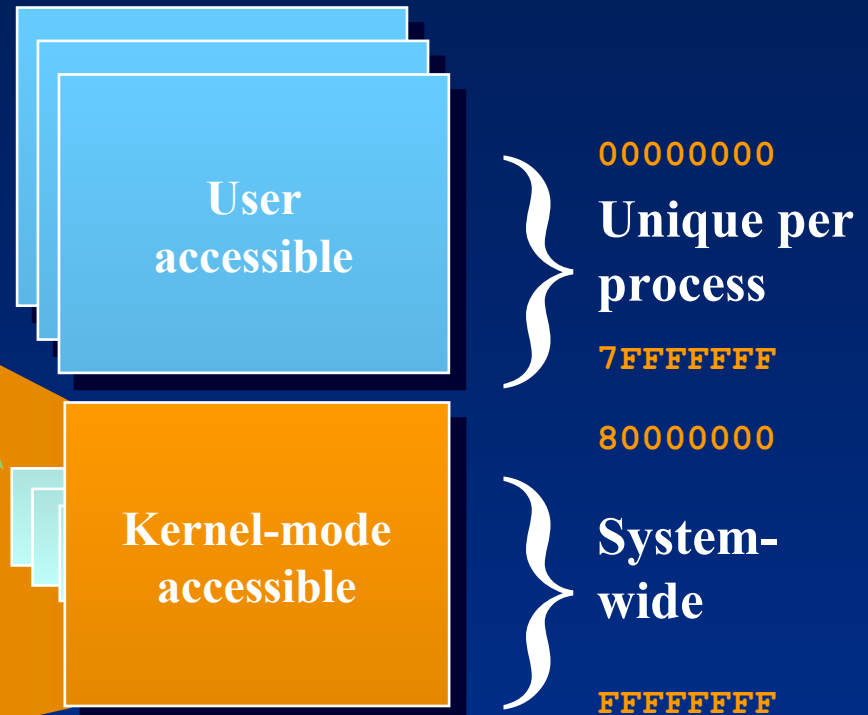
- The application you're running (.EXE and .DLLs)
- A user-mode stack for each thread (automatic storage)
- All static storage defined by the application



Virtual Address Space (V.A.S.)

System space contains:

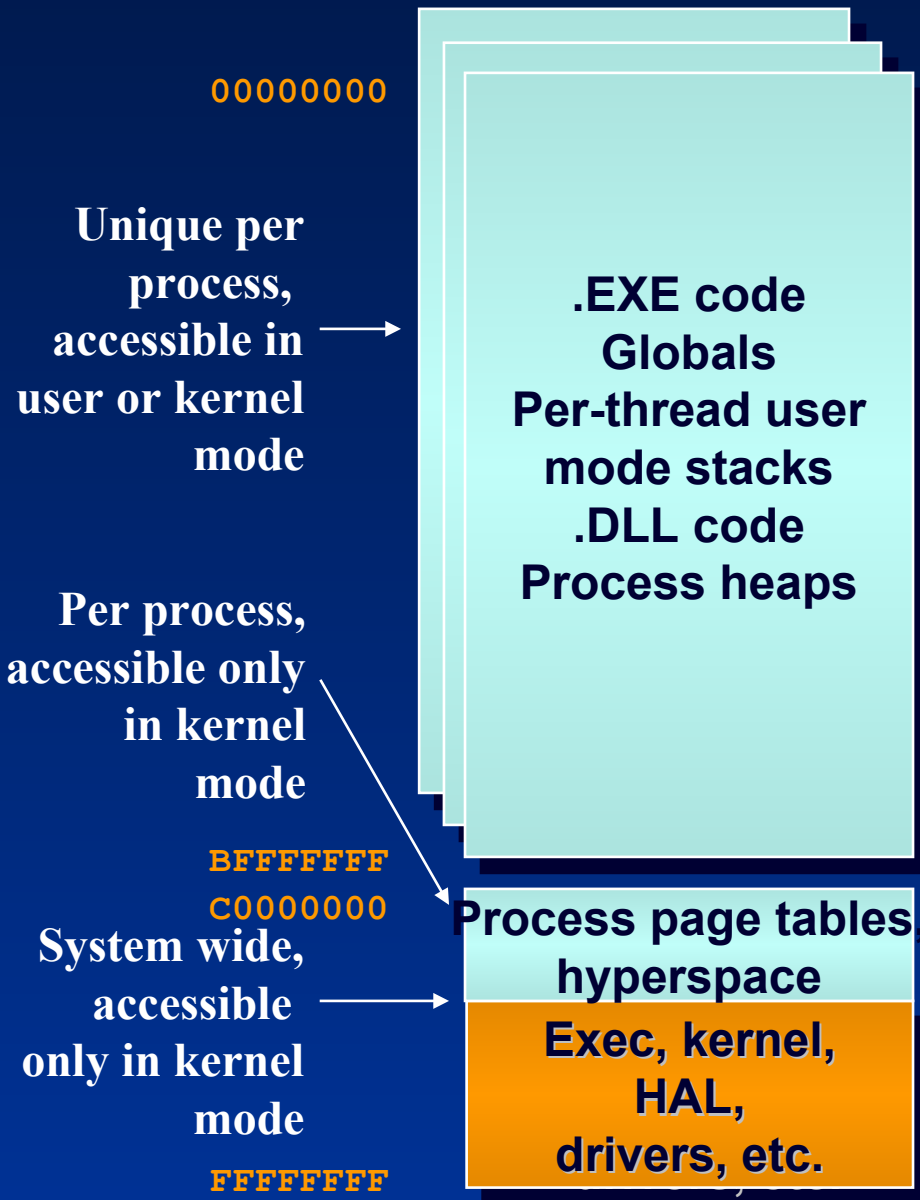
- Executive, kernel, and HAL
- Statically-allocated system-wide data cells
- Page tables (remapped for each process)
- Executive heaps (pools)
- Kernel-mode device drivers (in nonpaged pool)
- File system cache
- A kernel-mode stack for every thread in every process



32-bit Windows User Process Address Space Layout

Range	Size	Function
0x0 – 0xFFFF	64 KB	No-access region to catch incorrect pointer ref.
0x10000 - 0x7FFEFFFF	2 GB minus at least 192kb	The private process address space
0x7FFDE000 - 0x7FFDEFFF	4 KB	Thread Environment Block (TEB) for first thread, more TEBs are created at the page prior to that page
0x7FFDF000 - 0x7FFDFFFF	4 KB	Process Environment Block (PEB)
0x7FFE0000 - 0x7FFE0FFF	4 KB	Shared user data page – read-only, mapped to system space, contains system time, clock tick count, version number (avoid kernel-mode transition)
0x7FFE1000 – 0x7FFEFFFF	60 KB	No-access region
0x7FFF0000 – 0x7FFFFFFF	64 KB	No-access region to prevent threads from passing buffers that straddle user/system space boundary

3GB Process Space Option



- Only available on:
 - Windows 2003 Server, Enterprise Edition & Windows 2000 Advanced Server, XP SP2
 - Limits physical memory to 16 GB
 - /3GB option in BOOT.INI
 - Windows Server 2003 and XP SP2 supports variations from 2GB to 3GB (/USERVA=)
- Provides 3 GB per-process address space
 - Commonly used by database servers (for file mapping)
 - .EXE must have "large address space aware" flag in image header, or they're limited to 2 GB (specify at link time or with imagecfg.exe from ResKit)
 - Chief "loser" in system space is file system cache
 - Better solution: address windowing extensions (AWE)
 - Even better: 64-bit Windows

Large Address Space Aware Images

- Images marked as “large address space aware”:
 - Lsass.exe – Security Server
 - Inetinfo.exe—Internet Information Server
 - Chkdsk.exe – Check Disk utility
 - Dllhst3g.exe – special version of Dllhost.exe (for COM+ applications)
 - Esentutl.exe - jet database repair tool

- To see this type:

```
Imagecfig \windows\system32\*.exe > large_images.txt
```

- Then search for “large” in large_images.txt

Large Address Space Aware on 64-bits

- Images marked large address space aware get a full 4 GB process virtual address space
 - OS isn't mapped there, so space is available for process



Windows Task Manager

File Options View Help

Applications Processes Performance Networking Users

Image Name	PID	User Name	CPU	Mem Usage	VM Size
leakyapp3gb.exe	1848	Administrator	10	2,588,184 K	4,126,104 K
LEAKYAPP.EXE	188	Administrator	00	512,040 K	2,058,224 K

Physical Memory

- Maximum on Windows NT 4.0 was 4 GB for x86 (8 GB for Alpha AXP)
 - This is fixed by page table entry (PTE) format
- What about x86 systems with > 4 GB?
 - Pentium Pro and Xeon systems can support up to 64 GB physical memory
 - Four more bits of physical address in PTEs = 36 bits = 64 GB
- NT 4.0: Intel provided a driver that allows the use of RAM beyond 4 GB as a RAM disk
- Windows 2000 added proper support for PAE
 - Requires booting /PAE to select the PAE kernel
- Actual physical memory usable varies by Windows package (*see next slide*)

Physical Memory Limits (in GB)

	x86 HW 32-bit OS	x64 HW 32-bit OS	x64 HW 64-bit OS	IA-64 HW 64-bit OS
XP Home	4	4	n/a	n/a
XP Professional	4	4	16	n/a
Server 2003 Web Edition	2	2	n/a	n/a
Server 2003 Standard	4	4	16	n/a
Server 2003 Enterprise	32	32	64	64
Server 2003 Datacenter	64	128	1024	1024


Physical Memory Usage on Systems in PAE Mode

Virtual address space is still 4 GB, so how can you “use” > 4 GB of memory?

1. Although each process can only address 2 GB, many may be in memory at the same time (e.g. $5 * 2 \text{ GB processes} = 10 \text{ GB}$)

2. Files in system cache remain in physical memory

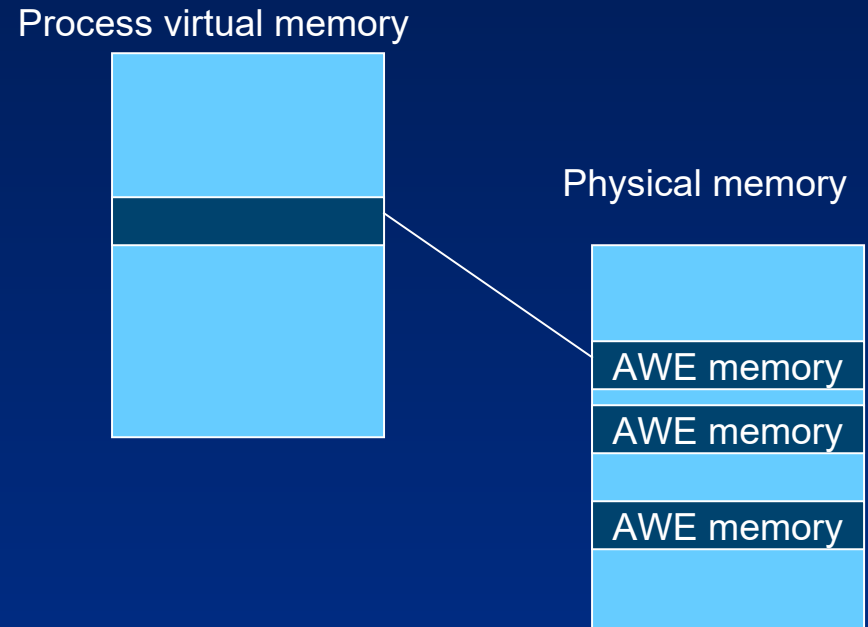
Although file cache doesn't know it, memory manager keeps unmapped data in physical memory



3. New Address Windowing Extensions (AWE) allow Windows processes to use more than 2 GB of memory

Address Windowing Extensions

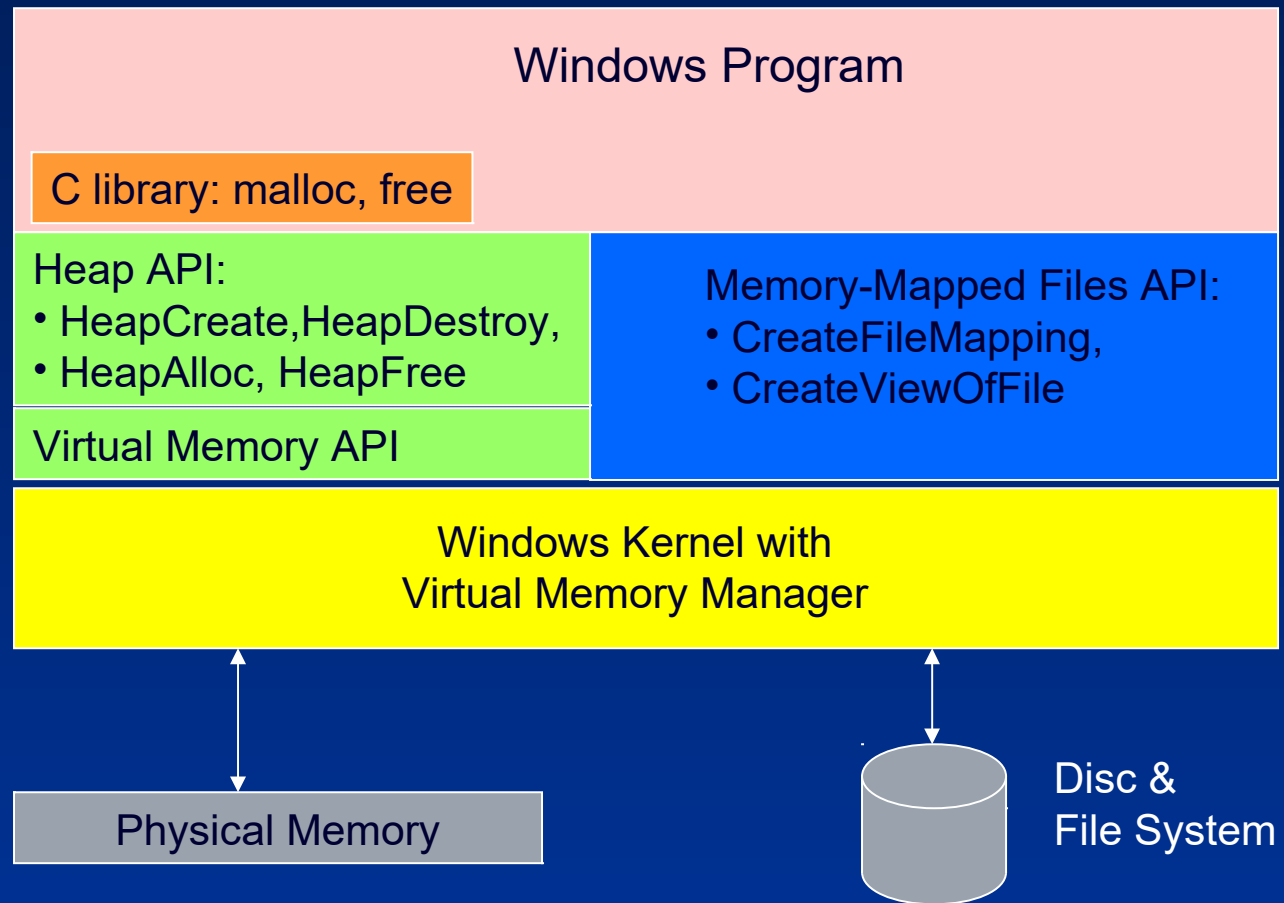
- AWE functions allow Windows processes to allocate large amounts of physical memory and then map “windows” into that memory
- Applications: database servers can cache large databases
- Up to programmer to control
 - Like DOS enhanced memory (EMS) with more bits...
- 64-bits removes this need



Windows Memory Allocation APIs

- HeapCreate, HeapAlloc, etc. (process heap APIs)
 - Windows equivalent of malloc(), free(), etc.
- VirtualAlloc(MEM_RESERVE)
- VirtualAlloc(MEM_COMMIT)
- VirtualFree
- VirtualQuery

Windows API Memory Management Architecture



Windows Memory Management

- Windows maintains pools of memory in heaps
- A process can contain several heaps – a default one and any number of private heaps (by explicit creation of them)
 - C library functions manage default heap: malloc, free, calloc
- Heaps are Windows objects – have handle
 - Each process has own default heap
 - Return value of NULL indicates failure (instead of INVALID_HANDLE_VALUE)

```
HANDLE GetProcessHeap( VOID );  
HANDLE HeapCreate (DWORD floptions,  
                  DWORD dwInitialSize,  
                  DWORD dwMaximumSize);  
BOOL HeapDestroy( HANDLE hHeap );
```

Managing Heap Memory

```
LPVOID HeapAlloc( HANDLE hHeap,  
                  DWORD dwFlags,  
                  DWORD dwBytes );
```

- dwFlags:
 - HEAP_GENERATE_EXCEPTION,
 - raise SEH on memory allocation failure
 - STATUS_NO_MEMORY, STATUS_ACCESS_VIOLATION
 - HEAP_NO_SERIALIZE:
 - no serialization of concurrent (multithreaded) requests
 - HEAP_ZERO_MEMORY: initialize allocated memory to zero
- dwSize:
 - Block of memory to allocate
 - For non-growable heaps: 0x7FFF8 (0.5 MB)
- HeapFree(), HeapReAlloc(),
- HeapCompact(), HeapValidate()

HeapLock(), HeapUnlock():
Manage concurrent accesses
to heap

Excerpt:

Sorting with Binary Search Tree

```
#define NODE_HEAP_ISIZE 0x8000
__try {
    /* Open the input file. */
    hIn = CreateFile (fname, GENERIC_READ, 0, NULL,
        OPEN_EXISTING, 0, NULL);
    if (hIn == INVALID_HANDLE_VALUE)
        fprintf(stderr, "Failed to open input file"), exit(1);
    /* Allocate the two heaps. */
    hNode = HeapCreate (
        HEAP_GENERATE_EXCEPTIONS | HEAP_NO_SERIALIZE,
        NODE_HEAP_ISIZE, 0);
    hData = HeapCreate (
        HEAP_GENERATE_EXCEPTIONS | HEAP_NO_SERIALIZE,
        DATA_HEAP_ISIZE, 0);
    /* Process the input file, creating the tree, actual search. */
    pRoot = FillTree (hIn, hNode, hData);
}
```


Heap Management Example (contd.)

```
/* Display the tree in Key order. */
printf ("Sorted file: %s"), fname); Scan (pRoot);


/* Destroy the two heaps and data structures. */
HeapDestroy (hNode); hNode = NULL;
HeapDestroy (hData); hData = NULL;
CloseHandle (hIn);
} /* End of main file processing and try block. */

__except (EXCEPTION_EXECUTE_HANDLER) {
    if (hNode != NULL) HeapDestroy (hNode);
    if (hData != NULL) HeapDestroy (hData);
    if (hIn != INVALID_HANDLE_VALUE) CloseHandle (hIn);
}
return 0;
```

- UNIX C library uses only a single heap
- UNIX sbrk() can create a Process' address space – no general-purpose MM
- UNIX does not generate signals on memory alloc.



Virtual Address Space Descriptors (VADs)




- VADs describe layout of virtual address space
 - Not the page mappings
- Used by memory manager to interpret access faults
 - Assists in "lazy evaluation"

See kernel debugger
command:
`!vad`



Example: Reserving Address Space

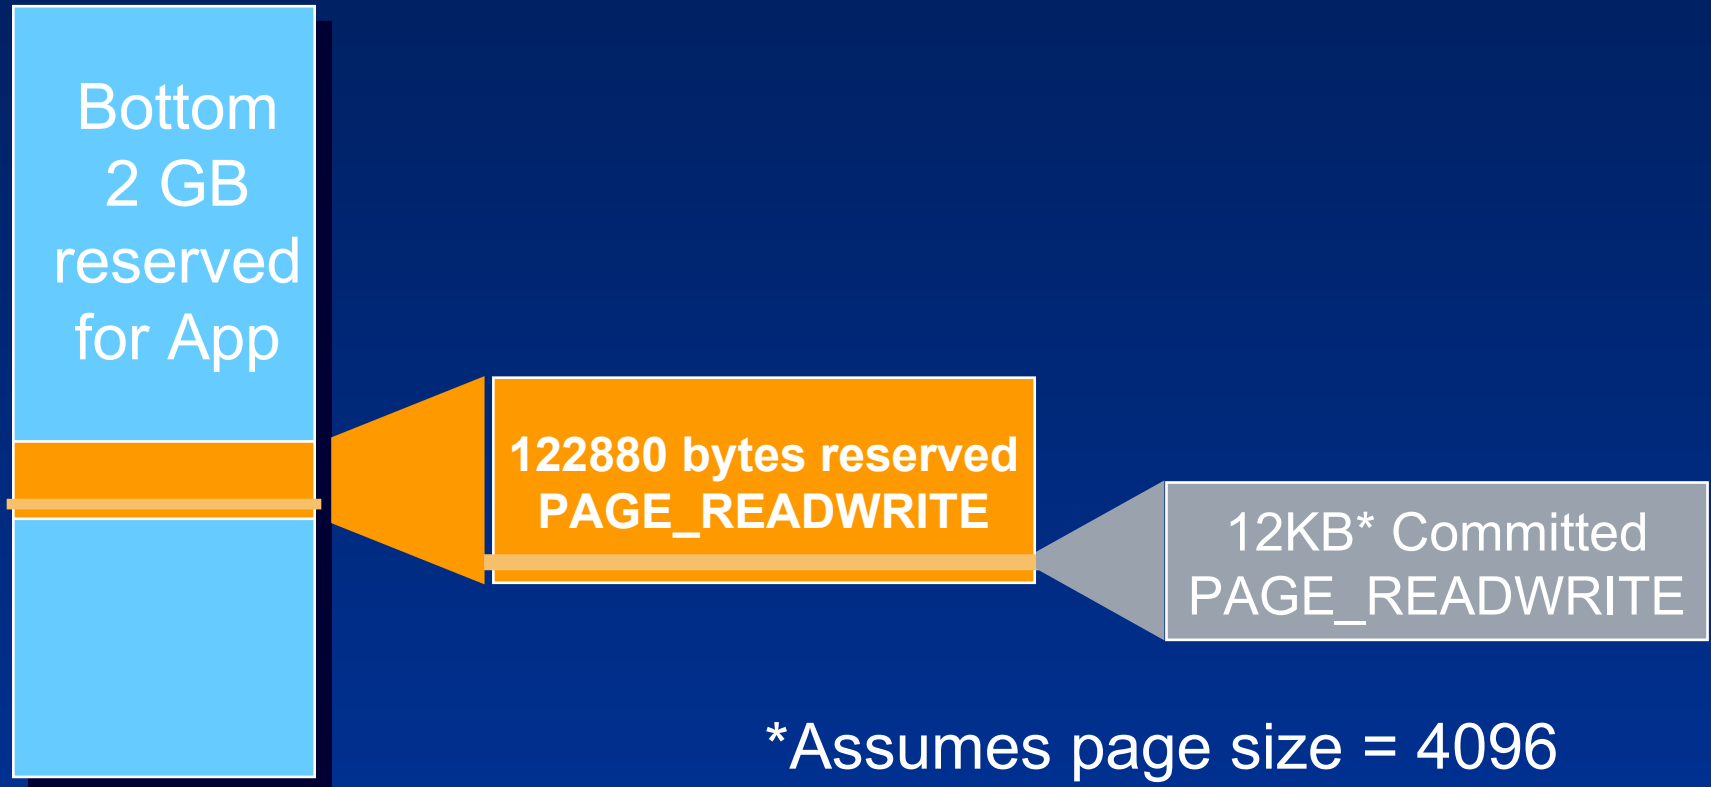
```
LPVOID lpMem = VirtualAlloc(NULL, 120000, MEM_RESERVE,  
PAGE_READWRITE);
```





Example: Committing Address Space

```
VirtualAlloc(lpMem + 6 * 1024, 7 * 1024, MEM_COMMIT,  
PAGE_READWRITE);
```



Memory-Mapped Files

- No need to perform direct file I/O (read/write)
- Data structures will be saved – be careful with pointers
- Convenient & efficient in-memory algorithms:
 - Can process data much larger than physical memory
- Improved performance for file processing
- No need to manage buffers and file data
 - OS does the hard work: efficient & reliable
- Multiple processes can share memory
- No need to consume space in paging file

File Mapping Object

```
HANDLE CreateFileMapping (HANDLE hFile,  
                           LPSECURITY_ATTRIBUTES lpsa,  
                           DWORD fdwProtect,  
                           DWORD dwMaximumSizeHigh,  
                           DWORD dwMaximumSizeLow,  
                           LPCTSTR lpszMapName );
```

Parameters:

- hFile:
 - hFile: handle to open file with compatible access rights (fdwProtect)
 - hFile == 0xFFFFFFFF: paging file, no need to create separate file
- fdwProtect:
 - PAGE_READONLY, PAGE_READWRITE, PAGE_WRITECOPY
- dwMaximumSizeHigh, dwMaximumSizeLow:
 - Zero: current file size is used
- lpszMapName:
 - Name of mapping object for sharing between processes or NULL

Shared Memory

```
HANDLE OpenFileMapping (HANDLE hFile,  
                        DWORD dwDesiredAccess,  
                        BOOL blInheritHandle,  
                        LPCTSTR lpName );
```

- Open an existing mapping object
 - Name comes from previous CreateFileMapping() call
 - First process creates mapping, subsequent processes open mapping
- dwDesiredAccess: same as fdwProtect
- lpName: name created with CreateFileMapping()
- CloseHandle() destroys mapping handles

Mapping Process Address Space to Mapping Objects

UNIX:
4.3BSD/SysV.4
have `mmap()` call

See also
`shmget()`, `shmctl()`,
`shmat()`, `shmdt()`

```
LPVOID MapViewOfFile( HANDLE hMapObject,  
                     DWORD fdwAccess, DWORD dwOffsetHigh,  
                     DWORD dwOffsetLow, DWORD cbMap );  
  
BOOL UnmapViewOfFile ( LPVOID lpBaseAddress );
```

- Allocate virtual memory space and map it to a file through a mapping object
 - Similar to HeapAlloc – much coarser granularity
 - Pointer to allocated block is returned (file view)
- Parameters:
 - `FILE_MAP_WRITE`, `FILE_MAP_READ`, `FILE_MAP_ALL_ACCESS` flag bits for `fdwAccess`
 - `cbMap`: size; entire file if zero
- `FlushViewOfFile()` : create consistent view

Limitation: 2GB virtual address space

Example: File Conversion with Memory Mapping (Excerpt)

```
/* Open the input file. */
hIn = CreateFile (fIn, GENERIC_READ, 0, NULL,
                 OPEN_EXISTING, FILE_ATTRIBUTE_NORMAL, NULL);
if (hIn == INVALID_HANDLE_VALUE) fprintf(stderr, "Failure opening input file."), exit(1);
/* Create a file mapping object on the input file. Use the file size. */
hInMap = CreateFileMapping (hIn, NULL, PAGE_READONLY, 0, 0, NULL);
if (hInMap == INVALID_HANDLE_VALUE) fprintf(stderr, "Failure Creating input map."), exit(2);

pInFile = MapViewOfFile (hInMap, FILE_MAP_READ, 0, 0, 0);
if (pInFile == NULL) fprintf(stderr, "Failure Mapping input file."), exit(3);
/* The output file MUST have Read/Write access for the mapping to succeed. */
hOut = CreateFile (fOut, GENERIC_READ | GENERIC_WRITE,
                 0, NULL, CREATE_ALWAYS, FILE_ATTRIBUTE_NORMAL, NULL);
if (hOut == INVALID_HANDLE_VALUE) fprintf(stderr, "Failure Opening output file."), exit(4);
hOutMap = CreateFileMapping (hOut, NULL, PAGE_READWRITE, 0, 2 * FsLow, NULL);
if (hOutMap == INVALID_HANDLE_VALUE) fprintf(stderr, "Failure creating output map."), exit(5);

pOutFile = MapViewOfFile (hOutMap, FILE_MAP_WRITE, 0, 0, 2 * FsLow);
if (pOutFile == NULL) fprintf(stderr, "Failure mapping output file."), exit(6);
```

Example (contd.)

```
pIn = pInFile;          /* actual file conversion */
pOut = pOutFile;
while (pIn < pInFile + FsLow) {
    *pOut = (WCHAR) *pIn; pIn++; pOut++;
}

/* Close all views and handles. */
UnmapViewOfFile (pOutFile); UnmapViewOfFile (pInFile);
CloseHandle (hOutMap); CloseHandle (hInMap);
CloseHandle (hIn); CloseHandle (hOut);
Complete = TRUE; return TRUE;
}

_except (EXCEPTION_EXECUTE_HANDLER) {
    /* Delete the output file if the operation did not complete successfully. */
    if (!Complete)
        DeleteFile (fOut);
    return FALSE;
}
```



Memory Management APIs

- Memory protection may be changed
 - per-page basis

```
status = VirtualProtect(baseAddress, size, newProtect, pOldprotect);
```

- Page protection choices:

PAGE_NOACCESS

PAGE_EXECUTE

PAGE_READONLY

PAGE_EXECUTE_READ

PAGE_READWRITE

PAGE_EXECUTE_READWRITE

PAGE_WRITECOPY

PAGE_EXECUTE_WRITECOPY

PAGE_GUARD

PAGE_NOCACHE



Memory Management Information

```
VOID GetSystemInfo(LPSYSTEM_INFO lpSystemInfo);
```

```
typedef struct _SYSTEM_INFO {  
    DWORD          dwOemId;  
    DWORD          dwPageSize;  
    LPVOID         lpMinimumApplicationAddress;  
    LPVOID         lpMaximumApplicationAddress;  
    DWORD          dwActiveProcessorMask;  
    DWORD          dwNumberOfProcessors;  
    DWORD          dwProcessorType;  
    DWORD          dwAllocationGranularity;  
    DWORD          dwReserved;  
} SYSTEM_INFO;
```



Querying Address Space

```
DWORD VirtualQuery( LPVOID lpAddress,  
                    PMEMORY_BASIC_INFORMATION lpBuffer, DWORD dwLength);
```

Returns:

```
typedef struct _MEMORY_BASIC_INFORMATION {  
    PVOID BaseAddress;           // Block base  
    PVOID AllocationBase;        // Region base  
    DWORD AllocationProtect;     // Region protection  
    DWORD RegionSize;           // # bytes in block  
    DWORD State;                 // State of block:  
                                // MEM_RESERVE, MEM_COMMIT, MEM_FREE  
    DWORD Protect;               // Pages protection  
    DWORD Type;                  // Type:  
                                // MEM_IMAGE, MEM_MAPPED, MEM_PRIVATE  
} MEMORY_BASIC_INFORMATION;
```



Memory Management Information

```
VOID GlobalMemoryStatus(LPMEMORYSTATUS lpms);
```

```
typedef struct _MEMORYSTATUS {  
    DWORD          dwLength;           // sizeof(MEMORYSTATUS)  
    DWORD          dwMemoryLoad;  
    DWORD          dwTotalPhys;  
    DWORD          dwAvailPhys;  
    DWORD          dwTotalPageFile;  
    DWORD          dwAvailPageFile;  
    DWORD          dwTotalVirtual;     // Process specific  
    DWORD          dwAvailVirtual;     // Process specific  
} MEMORYSTATUS, *LPMEMORYSTATUS;
```

Note: much more available via Registry Performance counters

Further Reading

- Pavel Yosifovich, Alex Ionescu, et al., “Windows Internals”, 7th Edition, Microsoft Press, 2017.
 - Chapter 5 – Memory management (from pp. 421)
 - Introduction to the memory manager (from pp. 421)
 - Services provided by the memory manager (from pp. 431)
- Jeffrey Richter, Programming Applications for Microsoft Windows, 4th Edition, Microsoft Press, September 1999.
 - Chapter 5 - Windows API Memory Architecture
 - Chapter 7 - Using Virtual Memory
 - Chapter 8 - Memory-Mapped Files
 - Chapter 9 - Heaps